

Property-Directed Inference of Universal Invariants or Proving Their Absence

ALEKSANDR KARBYSHEV, Aarhus University
NIKOLAJ BJØRNER, Microsoft Research
SHACHAR ITZHAKY, Massachusetts Institute of Technology
NOAM RINETZKY and SHARON SHOHAM, Tel Aviv University

We present *Universal Property Directed Reachability* (PDR^\forall), a property-directed semi-algorithm for automatic inference of invariants in a universal fragment of first-order logic. PDR^\forall is an extension of Bradley's $\text{PDR}/\text{IC3}$ algorithm for inference of propositional invariants. PDR^\forall terminates when it discovers a concrete counterexample, infers an inductive universal invariant strong enough to establish the desired safety property, or finds a *proof that such an invariant does not exist*. PDR^\forall is not guaranteed to terminate. However, we prove that under certain conditions, for example, when reasoning about programs manipulating singly linked lists, it does.

We implemented an analyzer based on PDR^\forall and applied it to a collection of list-manipulating programs. Our analyzer was able to automatically infer universal invariants strong enough to establish memory safety and certain functional correctness properties, show the absence of such invariants for certain natural programs and specifications, and detect bugs. All this without the need for user-supplied abstraction predicates.

CCS Concepts: • **Theory of computation** → **Invariants**; **Program verification**; *Verification by model checking*;

Additional Key Words and Phrases: Universal invariants, property-directed reachability, IC3, PDR, EPR

ACM Reference Format:

Aleksandr Karbyshev, Nikolaj Bjørner, Shachar Itzhaky, Noam Rinetzky, and Sharon Shoham. 2017. Property-directed inference of universal invariants or proving their absence. *J. ACM* 64, 1, Article 7 (March 2017), 33 pages.
DOI: <http://dx.doi.org/10.1145/3022187>

1. INTRODUCTION

We present *Universal Property Directed Reachability* (PDR^\forall), a semi-algorithm for automatic inference of quantified inductive invariants, and its application for the analysis of programs that manipulate unbounded data structures such as singly linked and doubly linked list data structures. For a correct program, the inductive invariant generated ensures that the program satisfies its specification. For an erroneous program, PDR^\forall produces a concrete counterexample. Historically, this has been addressed

This work was supported by EU FP7 project ADVENT (308830), ERC Grant No. [321174-VSSC], by Broadcom Foundation and Tel Aviv University Authentication Initiative, and by BSF Grant No. 2012259.

Authors' addresses: A. Karbyshev, Department of Computer Science, Aarhus University, Aabogade 34, DK-8200, Aarhus N, Denmark; email: karbyshev@cs.au.dk; N. Bjørner, Microsoft Research, One Microsoft Way, Washington, 98052; email: nbjorner@microsoft.com; S. Itzhaky, Massachusetts Institute of Technology, CSAIL, 32-G714, 32 Vassar Street, Cambridge, MA 02142; email: shachari@mit.edu; N. Rinetzky, School of Computer Science, Tel Aviv University, Tel Aviv 69978, Israel; email: maon@cs.tau.ac.il; S. Shoham, School of Computer Science, Tel Aviv University, Tel Aviv 69978, Israel; email: sharon.shoham@cs.tau.ac.il.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.

© 2017 ACM 0004-5411/2017/03-ART7 \$15.00

DOI: <http://dx.doi.org/10.1145/3022187>

by abstract interpretation [Cousot and Cousot 1977] algorithms, which automatically infer sound-inductive invariants, and bounded model checking algorithms, which explore a limited number of loop iterations in order to systematically look for bugs [Biere et al. 1999; Clarke et al. 2003]. We continue the line of recent work [Itzhaky et al. 2014; Albarghouthi et al. 2015] that simultaneously searches for invariants and counterexamples. We follow Bradley’s PDR/IC3 (Incremental Construction of Inductive Clauses for Indubitable Correctness) algorithm [Bradley 2011] by repeatedly strengthening a candidate invariant until it either becomes inductive or a counterexample is found.

In our experience, the correctness of many programs can be proven using universal invariants. Hence, we simplify matters by focusing on inferring universal first-order invariants. When PDR^\forall terminates, it yields one of the following outcomes: (i) a universal inductive invariant strong enough to show that the program respects the property, (ii) a concrete counterexample that shows that the program violates the desired safety property, or (iii) a *proof that the program cannot be proven correct using a universal invariant* in a given vocabulary.

Diagram Based Abstraction. Unlike previous work [Itzhaky et al. 2014; Albarghouthi et al. 2015], we neither assume that the predicates that constitute the invariants are known nor *a priori* bound the number of universal quantifiers. Instead, we rely on first-order theories with a *finite model property*: For such theories, SMT-based tools are able to either return UNSAT, indicating that the negation of a formula φ is valid, or construct a *finite model* σ of φ . We then translate σ into a *diagram* [Chang and Keisler 1990]—a formula describing the set of models that extend σ —and use the diagram to construct a *universal clause* to strengthen a candidate invariant.

Property-Directed Invariant Inference. Similarly to IC3, PDR^\forall iteratively constructs an increasing sequence of candidate inductive invariants F_0, \dots, F_N . Every F_i over-approximates the set \mathcal{R}_i of states that can be reached by up to i execution steps from a given set of *initial* states. In every iteration, PDR^\forall uses SMT to check whether one of the candidate invariants became inductive. If so, then the program respects the desired property. If not, then PDR^\forall iteratively strengthens the candidate invariants and adds new ones, guided by the considered property. Specifically, it checks if there exists a *bad* state σ that satisfies F_N but not the property. If so, then we use SMT again to check whether there is a state σ_a in F_{N-1} that can lead to a state in the *diagram* φ of σ in one execution step. If no such state exists, then the candidate invariant F_N can be strengthened by conjoining it with the negation of φ . Otherwise, we recursively strengthen F_{i-1} to exclude σ_a from its over-approximation of \mathcal{R}_{i-1} . If the recursive process tries to strengthen F_0 , then we stop and use a bounded model checker to look for a counterexample of length N . If no counterexample is found, then PDR^\forall determines that no universal invariant strong enough to prove the desired property exists (see Lemma 4.5). We note that PDR^\forall is not guaranteed to terminate. In Section 6, we show that under certain conditions, for example, when reasoning about programs manipulating singly linked lists, it does. Furthermore, in our experiments it terminates even when these conditions do not hold.

Example 1.1. Procedure `split()`, shown in Figure 1(a), moves the elements not satisfying the condition *ok* from the list pointed to by *h* to the list pointed to by *g*. PDR^\forall can infer tricky inductive invariants strong enough to prove several interesting properties: (i) memory safety, that is, no null dereference and no memory leaks; (ii) all the elements satisfying *ok* are kept in *h*; (iii) all the elements that do not satisfy *ok* are moved to *g*; (iv) no new elements are introduced; and (v) stability, that is, the reachability order between the elements satisfying *ok* is not changed. Our implementation verified that `split()` satisfies all the above properties fully automatically by inferring an inductive

```

void split(List h, List g) {
  i:=h;
  j:=null; k:=null;
  while (i≠null) {
    if ¬ok(i) then {
      if i=h then
        h:=i.n
      else
        j.n:=i.n;
      if g=null then
        g:=i
      else
        k.n:=i;
        k:=i; i:=i.n;
        k.n:=null
    }
    else { j:=i; i:=i.n }
  }
}

```

requires:

$$g = \text{null} \wedge \underline{h} = h \wedge (\forall x, y. n^*(x, y) \leftrightarrow \underline{n}^*(x, y))$$

ensures:

$$\begin{aligned}
 &(\forall z. h \neq \text{null} \wedge n^*(h, z) \rightarrow \text{ok}(z)) \wedge \\
 &(\forall z. g \neq \text{null} \wedge n^*(g, z) \rightarrow \neg \text{ok}(z)) \wedge \\
 &(\forall z. z \neq \text{null} \rightarrow (\underline{n}^*(h, z) \leftrightarrow n^*(h, z) \vee n^*(g, z))) \wedge \\
 &(\forall x, y. \\
 &\quad \underline{n}^*(h, x) \wedge \underline{n}^*(x, y) \wedge \text{ok}(x) \wedge \text{ok}(y) \rightarrow n^*(x, y)) \wedge \\
 &(\forall x, y. \\
 &\quad \underline{n}^*(h, x) \wedge \underline{n}^*(x, y) \wedge \neg \text{ok}(x) \wedge \neg \text{ok}(y) \rightarrow n^*(x, y))
 \end{aligned}$$

(a) A procedure that moves all the elements not satisfying $\text{ok}(\cdot)$ from list h to list g and its specification using pre- and post-conditions. Variables h, g, i, j , and k are pointers to list nodes, and $l.n$ denotes the “next” field of node l . $n^*(x, y)$ means a (possibly empty) path of n -fields from x to y . The ghost variables \underline{h} and $\underline{n}^*(\cdot, \cdot)$ record the head of the original list and the reachability order between its elements.

```

void filter(List h) {
  i:=h; j:=null;
  while (i≠null) {
    if ¬ok(i) then
      if i=h then
        h:=i.n
      else
        j.n:=i.n
      else j:=i;
      i:=i.n
    }
  }
}

```

requires:

true

ensures:

$$\forall z. h \neq \text{null} \wedge n^*(h, z) \rightarrow \text{ok}(z)$$

(b) A procedure that deletes all the nodes not satisfying $\text{ok}(\cdot)$ from list h .

Fig. 1. Motivating examples.

loop invariant consisting of 33 clauses (among them 17 are universal formulae). (The invariant is given in Appendix A.)

Example 1.2. Procedure `filter()`, shown in Figure 1(b), removes and deallocates the elements not satisfying the condition ok from the list pointed to by h . Figure 2 shows the loop invariant inferred by PDR^\forall when it was asked to verify a simplified version of property (iii): All the elements that do not satisfy ok are removed from h . The invariant highlights certain interesting properties of `filter()`. For example, clause L_4 says that if the head element of the list was processed and kept in the list (this is the only way $i \neq h$ can hold), then j becomes an immediate predecessor of i . Clause L_7 says that all the elements x reachable from h and not satisfying ok must occur after j .

Experimental Evaluation. We implemented PDR^\forall on top of the decision procedure of Itzhaky et al. [2014], and applied it to a collection of procedures that manipulate (possibly sorted) singly linked lists, doubly linked lists, and multi-linked lists. Our analysis

$$\begin{aligned}
I &= L_1 \wedge L_2 \wedge L_3 \wedge L_4 \wedge L_5 \wedge L_6 \wedge L_7 \quad \text{where} \\
L_1 &= i \neq h \wedge i \neq \text{null} \rightarrow n^*(j, i) \\
L_2 &= i \neq h \rightarrow \text{ok}(h) \\
L_3 &= n^*(h, j) \vee i \neq j \\
L_4 &= \forall x. i \neq h \wedge n^*(j, x) \wedge x \neq j \rightarrow n^*(i, x) \\
L_5 &= i \neq h \rightarrow \text{ok}(j) \\
L_6 &= \forall x. x = h \vee j = \text{null} \vee \neg n^*(h, x) \vee n^*(h, j) \vee \neg \text{ok}(j) \\
L_7 &= \forall x. j \neq \text{null} \wedge n^*(h, x) \wedge x \neq h \wedge \neg \text{ok}(x) \rightarrow n^*(j, x)
\end{aligned}$$

Fig. 2. Invariant for `filter()`. For better readability, some of the inferred clauses are displayed in the form of implications.

successfully verified interesting specifications, detected bugs in incorrect programs, and established the absence of universal invariants for certain correct programs.

Main Contributions. The main contributions of this work can be summarized as follows.

- We present PDR^\forall , a pleasantly simple, yet surprisingly powerful, combination of PDR [Bradley 2011] with a strengthening technique based on diagrams [Chang and Keisler 1990]. PDR^\forall enjoys a high degree of automation, because it does *not* require predefined abstraction predicates.
- The diagram-based abstraction is particularly interesting as it is determined “on-the-fly” according to the structural properties of the bad states discovered in PDR’s traversal of the state space.
- We prove that the diagram-based abstraction is precise in the sense that if PDR^\forall finds a spurious counterexample, then the program cannot be proven correct using a universal invariant. We believe that this is a unique feature of our approach.
- We provide sufficient conditions that ensure that PDR^\forall terminates.
- We implemented PDR^\forall on top of a decision procedure for the logic EA^R [Itzhaky et al. 2013]¹ and applied it successfully to verify a collection of list-manipulating programs, detect bugs, and prove the absence of universal invariants. We show that our technique outperforms an existing state-of-the-art PDR-based verification technique [Itzhaky et al. 2014] that uses the same decision procedure but requires user-supplied abstraction predicates. The implementation is available for download at <https://bitbucket.org/tausigplan/updr-distrib/>.
- The modeling of acyclic lists is based on the encoding developed in Itzhaky et al. [2013]. We also present a novel encoding that allows to model programs that manipulate (restricted) *cyclic* lists in EA^R and to apply our analysis to them.

2. PRELIMINARIES

This section formalizes the verification problem and sets terminology and notation. We start by explaining the way in which we use first-order logic to represent a transition system, which consists of a set of states and transitions between states. We then explain how we translate a program into a transition system and obtain a verification problem that captures the correctness of the program.

¹In Itzhaky et al. [2013], the logic AE^R was presented, whose *validity* is decidable. In this article, we are interested in *satisfiability* and consider the logic EA^R . The negation of an EA^R -formula is an AE^R -formula, and, hence, (un)satisfiability of EA^R can be reduced to validity of AE^R and is hence decidable.

2.1. Verification Problems and Their Representation in First-Order Logic

States. A state is represented by a *finite*² first-order model $\sigma = (D, \mathcal{I})$ over a vocabulary \mathcal{V} that consists of constants and relation symbols, where D is the finite *domain* of the model, and \mathcal{I} is the interpretation function of the symbols in \mathcal{V} . We assume that the domain D of every state is a subset of a fixed set \mathcal{U} , called a *universe*.

Transition Relation. The set of transitions of a transition system is defined using a *transition relation*. A transition relation is a set of models of a *double vocabulary* $\hat{\mathcal{V}} = \mathcal{V} \uplus \mathcal{V}'$, where vocabulary \mathcal{V} is used to describe the *source* state of the transition and vocabulary $\mathcal{V}' = \{v' \mid v \in \mathcal{V}\}$ is used to describe its *target* state: A model $\sigma' = (D, \mathcal{I}')$ over \mathcal{V}' describes a program state $\sigma = (D, \mathcal{I})$, where $\mathcal{I}(v) = \mathcal{I}'(v')$ for every symbol $v \in \mathcal{V}$.

Definition 2.1 (Reduct). Let $\hat{\sigma} = (D, \mathcal{I})$ be a model of $\hat{\mathcal{V}}$, and let $\Sigma \subseteq \hat{\mathcal{V}}$. The *reduct* of $\hat{\sigma}$ to Σ , denoted $\text{reduct}_{\Sigma}(\hat{\sigma})$, is the model $(D, \mathcal{I}_{\Sigma})$ of Σ where for every symbol $v \in \Sigma$, $\mathcal{I}_{\Sigma}(v) = \mathcal{I}(v)$.

We often write a transition $\hat{\sigma}$ as a pair of states (σ_1, σ_2) , such that σ_1 is the *reduct* of $\hat{\sigma}$ to vocabulary \mathcal{V} , and σ_2 is the state described by the *reduct* to \mathcal{V}' . We say that σ_2 is a successor of σ_1 , and σ_1 is a predecessor of σ_2 .

Verification Problem. A *transition system* is represented by a pair $TS = (\text{Init}, \rho)$, where Init is a closed first-order formula over \mathcal{V} used to denote the *initial states* of the program, and ρ is a closed formula over $\hat{\mathcal{V}}$ used to denote its *transition relation*. A state σ is *initial* if $\sigma \models \text{Init}$, and a pair of states (σ_1, σ_2) is a transition if $(\sigma_1, \sigma_2) \models \rho$. We say that a state is *reachable by at most i steps* of ρ (or *i -reachable* for short, when ρ is clear from the context) if it can be reached by at most i applications of ρ starting from some initial state. We denote the set of i -reachable states by \mathcal{R}_i . We say that a state is *reachable* if it is i -reachable for some i . We say that TS *satisfies* a *safety property* \mathcal{P} if all reachable states satisfy \mathcal{P} . We often define $\text{Bad} \stackrel{\text{def}}{=} \neg \mathcal{P}$ and refer to states satisfying Bad as *bad states*.

Properties and Assertions. *Properties* are sets of states. We express properties, such as pre- and post-conditions, and assertions within the loop body, using closed logical formulae over \mathcal{V} .

Invariants. An *invariant* of a transition system is a property that should hold for all reachable states. It is *inductive* if it is closed under application of ρ . In the following, we use $(\varphi)'$ to denote the formula obtained by replacing every constant and relation symbol in formula φ with its primed version.

Definition 2.2 (Invariants). Let $TS = (\text{Init}, \rho)$ be a transition system and \mathcal{P} a safety property over \mathcal{V} . A closed formula \mathcal{I} is a *safety inductive invariant* (*invariant*, in short) for TS and \mathcal{P} if (i) $\text{Init} \Rightarrow \mathcal{I}$, (ii) $\mathcal{I} \wedge \rho \Rightarrow (\mathcal{I})'$, and (iii) $\mathcal{I} \Rightarrow \mathcal{P}$.

If there exists an invariant for TS and \mathcal{P} , then TS satisfies \mathcal{P} . An invariant is *universal* if it is equivalent to a universal formula (i.e., a formula with a \forall^* quantifier prefix in prenex normal form). We note that the invariants inferred by PDR^{\forall} are conjunctions of *universal clauses*, where a universal clause is a universally quantified disjunction of literals (positive or negative atomic formulae).

2.2. From Programs to Verification Problems

Programs. We handle single loop programs, that is, we assume that a program has the form `while Cond do Cmd`, where `Cmd` is loop free. We encode more complicated

²All first-order models considered in this work are finite, that is, have a finite domain.

control structures, for example, nested or multiple loops, by explicitly recording the program counter. For clarity, in our examples we allow for a sequence of instructions preceding the loop. Technically, we encode their effect in the loop's precondition.

Program Semantics and Verification Problem. The semantics of a program is described by a transition system. We consider the states of the program at the beginning of each iteration of the loop. Each transition (σ_1, σ_2) describes one possible execution of the loop body, Cmd , that is, it relates the state σ_1 at the beginning of an iteration of the loop to the state σ_2 at the end of the iteration.

Technically, following Itzhaky et al. [2013], we derive the semantics of the loop body as a transition relation formula ρ from a *weakest liberal precondition* predicate transformer, wlp , defined for each command type. As an example, the top of Table II presents the definition of wlp for the simple imperative language IMP [Winskel 1993]. To construct the transition relation using wlp , we define an *identity* formula Id that specifies that the input and the output states are identical. That is, Id is a two-vocabulary closed formula such that $(\sigma, \sigma') \models Id \Leftrightarrow \sigma = \sigma'$. Formally, Id is defined by

$$Id \stackrel{\text{def}}{=} \bigwedge_{c \in \mathcal{C}} c = c' \wedge \bigwedge_{R \in \mathcal{R}} \forall \bar{\alpha}. R(\bar{\alpha}) \Leftrightarrow R'(\bar{\alpha}), \quad (1)$$

where \mathcal{C} and \mathcal{R} denote the sets of constants and relation symbols in \mathcal{V} , respectively, and $\bar{\alpha}$ is a list of variables according to the arity of the relation symbol R . The vocabulary \mathcal{V} corresponds to the structure σ , and \mathcal{V}' corresponds to σ' .

We then define

$$\rho \stackrel{\text{def}}{=} Cond \wedge wlp\llbracket Cmd \rrbracket Id, \quad (2)$$

where $wlp\llbracket Cmd \rrbracket$ denotes the weakest liberal precondition of the loop body.

We define $Init$ and Bad using the programs pre- and post-conditions, as well as its assertions:

$$Init \stackrel{\text{def}}{=} Pre \quad \text{and} \quad Bad \stackrel{\text{def}}{=} (\neg Cond \wedge \neg Post) \vee (Cond \wedge \neg wlp\llbracket Cmd \rrbracket \mathbf{true}). \quad (3)$$

That is, a state is initial if it satisfies the precondition, and it is bad in one of two cases: (i) if the state satisfies the negation of the loop condition (which indicates termination of the loop) but does not satisfy the post-condition (this captures the requirement that when the loop terminates the post-condition needs to hold) and (ii) if the state leads to a violation of an assertion within the loop body when it is encountered in the loop head. This is captured by the subformula $\neg wlp\llbracket Cmd \rrbracket \mathbf{true}$ of Bad .

The construction of ρ , $Init$, and Bad ensures that $TS = (Init, \rho)$ satisfies $\mathcal{P} = \neg Bad$ if and only if any execution of the program starting at a state that satisfies the given precondition never violates an assertion, and if it terminates, then it ends in a state that satisfies the postcondition.

3. REASONING ABOUT HEAP-MANIPULATING PROGRAMS USING EFFECTIVELY PROPOSITIONAL LOGIC

In this section, we exemplify how we represent heap-manipulating programs, such as the ones used in our running examples and experiments, as well as the corresponding verification problems in first-order logic.

We start by defining the fragment of logic used and continue to describe the construction of the formulae ρ , $Init$, and Bad for a program. First, we present the construction of the formulae for programs that manipulate *acyclic* data structures, as developed

Table I. Effectively Propositional Axiomatization of Deterministic Reflexive-Transitive Closure

$\Gamma_{n^*} \stackrel{\text{def}}{=} \forall \alpha, \beta. n^*(\alpha, \beta) \wedge n^*(\beta, \alpha) \leftrightarrow \alpha = \beta$	reflexivity + acyclicity
$\wedge \forall \alpha, \beta, \gamma. n^*(\alpha, \beta) \wedge n^*(\beta, \gamma) \rightarrow n^*(\alpha, \gamma)$	transitivity
$\wedge \forall \alpha, \beta, \gamma. n^*(\alpha, \beta) \wedge n^*(\alpha, \gamma) \rightarrow n^*(\beta, \gamma) \vee n^*(\gamma, \beta)$	semi-linearity

in Itzhaky et al. [2013]. Next, we develop a novel construction that also handles restricted *cyclic* data structures.

EPR and EA^R . Effectively-Propositional logic (*EPR*), also known as the Bernays-Schönfinkel-Ramsey class, is a fragment of first-order logic that allows for relational first-order formulae with a quantifier prefix of the form $\exists^* \forall^*$ but forbids functional symbols. Satisfiability of *EPR* is decidable. *EPR* enjoys the *small model property*: Every satisfiable formula in *EPR* is guaranteed to have a finite model [Lewis 1980].

In our running examples and experiments, we represent programs and the corresponding verification problems using EA^R [Itzhaky et al. 2013], an auxiliary logic built on top of *EPR*, which enables natural reasoning about programs manipulating linked-data structures. EA^R extends *EPR* by allowing a deterministic transitive-closure operator $*$ over acyclic relations. Satisfiability of EA^R is reducible to that of *EPR* and enjoys the same properties. Technically, the reduction introduces first-order axioms (*EPR* formulae) that provide a complete characterization of $*$. These axioms are given in Table I.

Programs Manipulating Linked-Data Structures as Transition Systems. To represent memory states of list manipulating programs, we fix an infinite countable universe \mathcal{U} whose individuals represent dynamically allocated objects. Recall that a state is represented by a finite first-order model $\sigma = (D, \mathcal{I})$ with $D \subseteq \mathcal{U}$.

We use a vocabulary \mathcal{V} that associates every program variable x with a constant x , contains a designated constant *null* to denote the null value, and contains the special binary predicate symbol $n^*(\cdot, \cdot)$ that defines reachability over every pointer field n , for example, in Example 1.1 and 1.2. Notice that n itself is not part of the vocabulary, but it is definable using the (open) formula φ_n in Table II. We use $n^+(\alpha, \beta)$ as a shorthand for $n^*(\alpha, \beta) \wedge \alpha \neq \beta$. Clearly, these definitions for φ_n and n^+ rely on acyclicity of n : If there was a cycle, then for every node u on the cycle we would have $n^+(u, u)$, and also for any two nodes u, v we would have $n^*(u, v)$, so there would not be enough information in n^* to define φ_n based on it. In particular, the order of the node in the cycle is not encoded in n^* .

In addition, we represent a Boolean function *ok* with a unary predicate $ok(\cdot)$ and an order relation (e.g., for sorting) with a binary predicate $R(\cdot, \cdot)$.

We depict memory states $\sigma = (D, \mathcal{I})$ as directed graphs (e.g., Figure 3). Individuals in D , representing heap locations, are depicted as circles labeled by their name. We draw an edge from the name of constant x or a unary predicate *ok* to an individual v if $\sigma \models x = v$ or $\sigma \models ok(v)$, respectively. For clarity, we do not directly depict the interpretation of the n^* relation. Instead, we use a more compact drawing scheme where we draw an n -annotated edge between v and u if $\sigma \models \varphi_n(v, u)$. The interpretation of n^* can be inferred from the n -annotated edges by (i) omitting the incoming edges of the element that corresponds to the *null* constant and (ii) considering the reflexive transitive closure of the remaining edges.

Transition Relation. We express the semantics of loop-free code as a transition relation ρ over the above vocabulary by defining a *weakest liberal precondition* predicate transformer, $wlp[\![-]\!]$, for each command type. We do this in a simple language IMP^R , which is an extension of *IMP* [Winskel 1993] with heap-related commands. The rules for wlp are shown in Table II. The notation $Q[t/x]$ is used to denote substitution of

Table II. Rules for Computing Weakest Liberal Preconditions for Procedures in IMP^R

$wlp\llbracket\text{skip}\rrbracket Q$	$\stackrel{\text{def}}{=} Q$
$wlp\llbracket x := y \rrbracket Q$	$\stackrel{\text{def}}{=} Q[y/x]$
$wlp\llbracket \text{Cmd}_1; \text{Cmd}_2 \rrbracket Q$	$\stackrel{\text{def}}{=} wlp\llbracket \text{Cmd}_1 \rrbracket (wlp\llbracket \text{Cmd}_2 \rrbracket Q)$
$wlp\llbracket \text{if } B \text{ then } \text{Cmd}_1 \text{ else } \text{Cmd}_2 \rrbracket Q$	$\stackrel{\text{def}}{=} \llbracket B \rrbracket \wedge wlp\llbracket \text{Cmd}_1 \rrbracket Q \vee \neg \llbracket B \rrbracket \wedge wlp\llbracket \text{Cmd}_2 \rrbracket Q$
$wlp\llbracket \text{assert } B \rrbracket Q$	$\stackrel{\text{def}}{=} \llbracket B \rrbracket \wedge Q$

$wlp\llbracket x.n := \text{null} \rrbracket Q$	$\stackrel{\text{def}}{=} Q[n^*(\alpha, \beta) \wedge (\neg n^*(\alpha, x) \vee n^*(\beta, x)) / n^*(\alpha, \beta)]$
$wlp\llbracket x.n := y \rrbracket Q$	$\stackrel{\text{def}}{=} \neg n^*(y, x) \wedge Q[n^*(\alpha, \beta) \vee (y \neq \text{null} \wedge n^*(\alpha, x) \wedge n^*(y, \beta)) / n^*(\alpha, \beta)]$
$wlp\llbracket x := y.n \rrbracket Q$	$\stackrel{\text{def}}{=} \forall \alpha. \varphi_n(y, \alpha) \rightarrow Q[\alpha/x]$
where $\varphi_n(s, t)$	$\stackrel{\text{def}}{=} (n^+(s, t) \wedge \forall \gamma. n^+(s, \gamma) \rightarrow n^*(t, \gamma)) \vee (t = \text{null} \wedge \forall \gamma. \neg n^+(s, \gamma))$
$wlp\llbracket x := \text{new} \rrbracket Q$	$\stackrel{\text{def}}{=} \forall \alpha. (\bigwedge_{p \in \text{vars} \cup \{\text{null}\}} \neg n^*(p, \alpha)) \rightarrow Q[\alpha/x]$

Q is a post-condition expressed as a first-order formula. The top frame shows the standard wlp rules for IMP, the bottom frame contains our additions for heap updates, dereference, and memory allocation. We assume that the program nullifies a field before modifying it, that is, every command of the form $x.n := y$ is preceded by a command $x.n := \text{null}$

Table III. Leading-existential Variant of wlp Rules

$wlp^\exists\llbracket x := y.n \rrbracket Q$	$\stackrel{\text{def}}{=} \exists \alpha. \varphi_n(y, \alpha) \wedge Q[\alpha/x]$
$wlp^\exists\llbracket x := \text{new} \rrbracket Q$	$\stackrel{\text{def}}{=} \exists \alpha. (\bigwedge_{p \in \text{vars} \cup \{\text{null}\}} \neg n^*(p, \alpha)) \wedge Q[\alpha/x]$

all the occurrences of the constant x in Q with the term t . The notation $Q[\varphi/n^*(\alpha, \beta)]$ denotes substitution of any atom of the form $n^*(\cdot, \cdot)$ in Q with the formula φ , where α and β may occur as term placeholders in φ and are filled in with the arguments of n^* . vars is the set of (constant symbols pertaining to) variables used in the program. As shown in Itzhaky et al. [2014], the rules for wlp are sound and complete.

The encoding of lists using n^* and the corresponding update rules wlp may seem confounding at first but follow a fairly simple intuition: When removing a pointer link, all paths that go through the changed node are disconnected; when adding a link, all paths into the source get connected to the (single) path from the target. This is expressed, respectively, by the formulae

$$n^*(\alpha, \beta) \wedge (\neg n^*(\alpha, x) \vee n^*(\beta, x)) \quad (\text{for } x.n := \text{null})$$

and

$$n^*(\alpha, \beta) \vee (y \neq \text{null} \wedge n^*(\alpha, x) \wedge n^*(y, \beta)) \quad (\text{for } x.n := y).$$

The former describes all n -paths except those that go through x , and the latter describes all n -paths with the addition of those that were connected by the new edge from x to y . Here, α and β denote arbitrary heap locations. When traversing a pointer ($x := y.n$), the successor can be expressed by its transitive closure n^* using the formula φ_n : For any two locations s and t (which are not null), the successor of s is t iff t is on the path starting at s (but not s itself), and no other node lies between s and t . The successor of s is null iff the path starting at s is empty. Our ability to recover the successor from n^* is key to having a complete encoding of heap structures using transitive closure.

Given the wlp definition, the transition relation ρ is defined as in Equation (2) (see Section 2.2). It is important to notice that, since φ_n is a universal formula occurring in a negative context in $wlp\llbracket x := y.n \rrbracket Q$, which is itself defined by a universal formula, the resulting formula ρ will have a quantifier prefix $\forall^*\exists^*$. We would like to get an EA^R formula for our purposes; we achieve this by changing the quantified rules slightly, resulting in the variants in Table III. All other rules remain unchanged. From here

onward, we switch to the definition of the transition relation as

$$\rho \stackrel{\text{def}}{=} \text{Cond} \wedge \text{wlp}^3 \llbracket \text{Cmd} \rrbracket \text{Id}. \quad (4)$$

Notice that the equivalence of semantics relies on the fact that for every location $s \in \mathcal{U}$ that differs from *null* there is exactly one location t such that $\varphi_n(s, t)$, as follows from the definition.

Example 3.1. Since the transition relation obtained for `filter()` is large, we demonstrate the construction of ρ on a simpler program, where the loop body consists of the following command:

$$\text{Cmd} = k := i.n; i.n := \text{null}; i := k.$$

We then have

$$\begin{aligned} \text{wlp}^3 \llbracket \text{Cmd} \rrbracket Q &= \text{wlp}^3 \llbracket k := i.n; i.n := \text{null}; i := k \rrbracket Q \\ &= \text{wlp}^3 \llbracket k := i.n \rrbracket (\text{wlp}^3 \llbracket i.n := \text{null} \rrbracket (\text{wlp}^3 \llbracket i := k \rrbracket Q)) \\ &= \exists \alpha. \varphi_n(i, \alpha) \wedge (\text{wlp}^3 \llbracket i.n := \text{null} \rrbracket (Q[k/i]))[\alpha/k] \\ &= \exists \alpha. \varphi_n(i, \alpha) \wedge ((Q[k/i])[n^*(\alpha, \beta) \wedge (\neg n^*(\alpha, i) \vee n^*(\beta, i))/n^*(\alpha, \beta)])[\alpha/k]. \end{aligned}$$

The transition relation is then constructed as follows:

$$\begin{aligned} \rho &= i \neq \text{null} \wedge \text{wlp}^3 \llbracket \text{Cmd} \rrbracket \text{Id} \\ &= i \neq \text{null} \wedge \text{wlp}^3 \llbracket \text{Cmd} \rrbracket (i = i' \wedge k = k' \wedge \forall \alpha, \beta. n^*(\alpha, \beta) \leftrightarrow n^*(\alpha, \beta)) \\ &= i \neq \text{null} \wedge \\ &\quad \exists \alpha. \varphi_n(i, \alpha) \wedge (\alpha = i' \wedge \alpha = k' \wedge \forall \alpha, \beta. (n^*(\alpha, \beta) \wedge (\neg n^*(\alpha, i) \vee n^*(\beta, i))) \leftrightarrow n^*(\alpha, \beta)). \end{aligned}$$

PROPOSITION 3.2. EA^R is closed under $\text{wlp}^3 \llbracket \text{Cmd} \rrbracket$; that is, if $Q \in EA^R$ then $\text{wlp}^3 \llbracket \text{Cmd} \rrbracket Q \in EA^R$. In particular, $\rho \in EA^R$ (as defined by Equation (4)).

Initial and Bad States. We express properties of list-manipulating programs, for example, their pre- and post-conditions, *Pre* and *Post*, respectively, using assertions written in EA^R over the above vocabulary. *Init* and *Bad* are defined based on these assertions, as shown in Equation (3) (see Section 2.2).

Example 3.3. In Example 1.2, we have $\text{Pre} = i = h \wedge j = \text{null}$ and $\text{Post} = h \neq \text{null} \rightarrow \forall z. n^*(h, z) \rightarrow \text{ok}(z)$. Note that these refer to the pre- and post-conditions that should hold right before the loop begins and right after it terminates, respectively. Therefore, $\text{Init} \stackrel{\text{def}}{=} i = h \wedge j = \text{null}$ and $\text{Bad} \stackrel{\text{def}}{=} i = \text{null} \wedge \neg(h \neq \text{null} \rightarrow \forall z. n^*(h, z) \rightarrow \text{ok}(z))$. Here, a state is bad if $i = \text{null}$ (i.e., it occurs when the loop terminates) and h points to a non-empty list that contains an element not having the property *ok*. In this example, there are no `assert` statements in the body, and, hence, the second disjunct in the definition of *Bad* in Equation (3), which captures the semantics of assertion violations, simplifies to **false**, and is subsumed by the first disjunct.

In our analysis, the EA^R formulae *Init*, *Bad*, and ρ are translated into equisatisfiable *EPR* formulae [Itzhaky et al. 2013].

3.1. Modeling of Programs Manipulating Cyclic Linked Lists

As an extension of previous work [Itzhaky et al. 2014; Karbyshev et al. 2015], which targeted acyclic data structures, we augment the formalism shown above to handle a restricted form of cycles. The new formalism allows at most one cycle to be present in the heap at any given time. This is achieved by decomposing the pointer edges, labeled

Table IV. Modified *wlp* Rules for Handling (Restricted) Cyclic Lists

$wlp[x.n := null] Q$	$\stackrel{\text{def}}{=} Q$	$\left[\begin{array}{l} \text{ite}(\psi_{km}(x, x), null, m_s)/m_s, \\ \text{ite}(\psi_{km}(x, x), null, m_t)/m_t, \\ k^*(\alpha, \beta) \wedge (\neg k^*(\alpha, x) \vee k^*(\beta, x)) \vee \\ (\psi_{km}(\alpha, \beta) \wedge \neg k^*(\alpha, x) \wedge k^*(\beta, x) \wedge x \neq m_s)/k^*(\alpha, \beta) \end{array} \right]$
where $\psi_{km}(s, t)$	$\stackrel{\text{def}}{=} m_s \neq null \wedge k^*(s, m_s) \wedge k^*(m_t, t)$	
$wlp[x.n := y] Q$	$\stackrel{\text{def}}{=} (k^*(y, x) \rightarrow m_s = null) \wedge$	Q
		$\left[\begin{array}{l} \text{ite}(k^*(y, x), x, m_s)/m_s, \\ \text{ite}(k^*(y, x), y, m_t)/m_t, \\ k^*(\alpha, \beta) \vee \\ (y \neq null \wedge \neg k^*(y, x) \wedge k^*(\alpha, x) \wedge k^*(y, \beta))/k^*(\alpha, \beta) \end{array} \right]$
$wlp[x := y.n] Q$	$\stackrel{\text{def}}{=} \forall \alpha. \varphi_n(y, \alpha) \rightarrow Q[\alpha/x]$	
where $\varphi_n(s, t)$	$\stackrel{\text{def}}{=} \begin{cases} t = m_t & s = m_s \\ (k^+(s, t) \wedge \forall \gamma. k^+(s, \gamma) \rightarrow k^*(t, \gamma)) \vee \\ (t = null \wedge \forall \gamma. \neg k^+(s, \gamma)) & \text{otherwise} \end{cases}$	

n , into a set of acyclic edges labeled k plus *at most one* additional edge labeled m . This is always possible—if the heap contains (at most) one cycle, then it is enough to remove (at most) one edge to make it acyclic.

We denote n^* and k^* as the reflexive transitive closures of n and k and $\langle m_s, m_t \rangle$ the source and destination of the edge labeled m , if it is present (if m is not present, $m_s = m_t = null$). The following relationship holds between n^* and k^* , m_s , m_t :

$$\forall \alpha, \beta. n^*(\alpha, \beta) \Leftrightarrow k^*(\alpha, \beta) \vee (m_s \neq null \wedge k^*(\alpha, m_s) \wedge k^*(m_t, \beta)). \quad (5)$$

Therefore, k^* , m_s , m_t fully characterize the heap reachability. The axioms in Table I now hold for k^* instead of n^* . In addition, we require that if m is present, then m_s has no k -successor, and the edge m closes the cycle; that is, $m_s \neq null \rightarrow k^*(m_t, m_s) \wedge \neg \exists \alpha. k^+(m_s, \alpha)$. We now modify the *wlp* formulae from Table II to reflect the new situation. The new semantics for $x.n := null$, $x.n := y$, and $x := y.n$ are shown in Table IV. The multiple substitutions in the brackets are done in parallel. The operator $\text{ite}(p, a, b)$ denotes a term that is equal to a if p is **true** and b otherwise.³ $\psi_{km}(s, t)$ is a formula expressing a path between s and t utilizing the special edge m . $\psi_{km}(u, u)$ means that u lies on the cycle.

The semantics maintains the edge m by creating it when a cycle is closed as a result of an assignment of the form $x.n := y$, and removing it or replacing it with a k edge when the cycle is broken by $x.n := null$. Notice that $wlp[x := y.n]$ remains as in Table II, except that the definition of φ_n is changed. The adjustment in Table III for wlp^\exists is suitable in this case as well.

4. UNIVERSAL-PROPERTY-DIRECTED REACHABILITY

In this section, we present *Universal Property Directed Reachability* (PDR^\forall), an algorithm for checking if a transition system TS satisfies a safety property \mathcal{P} . PDR^\forall is an adaptation of Bradley's *property-directed reachability* (IC3) algorithm [Bradley 2011]

³Any formula containing ite can be translated to an equivalent first-order formula using only standard connectives; however, SMT-LIB-compliant solvers natively support ite [Barrett et al. 2010], so this translation is not needed.

that uses universal formulae instead of propositional predicates [Bradley 2011; Eén et al. 2011; Hoder and Bjørner 2012] or predicate abstraction [Itzhaky et al. 2014]. We use Example 1.2 as a running example throughout this section.

We note that several flavors of PDR have been developed in the past few years, including both particular implementations [Bradley 2011; Eén et al. 2011] and abstract presentations [Hoder and Bjørner 2012]. In this article, we use a deterministic implementation of PDR that employs recursion instead of the more commonly used *obligation queue*. This allows us to simplify the presentation and focus on the unique aspects of PDR^\forall at the expense of precluding discussion of several optimizations.

Requirements. We require that the transition relation ρ , as well as the *Init* and *Bad* conditions, are expressible in a logic \mathcal{L} , which is a fragment of first-order logic. (We can partly handle transitive closure using the approach of Itzhaky et al. [2013]. See Section 3.) We assume that \mathcal{L} is closed under conjunction and contains all universal and existential formulae. We require that every satisfiable formula in \mathcal{L} has a finite model, and assume to have a decision procedure $\text{SAT}(\psi)$, which checks if a formula ψ in \mathcal{L} is satisfiable, and a function $\text{model}(\psi)$, which returns a *finite* model σ of ψ if such a model exists and *None* otherwise. *EPR* (and EA^R) is an example of a logic \mathcal{L} that can be used to express transition systems and properties. In particular, as \mathcal{P} is the negation of *Bad*, it has to be described by a negation of an *EPR* (or EA^R) formula, that is \mathcal{P} is a $\forall^*\exists^*$ formula.

4.1. Diagrams as Structural Abstractions

PDR^\forall iteratively strengthens a candidate invariant by retrieving program states that lead to bad states and checking whether the retrieved states are reachable. In that sense, PDR^\forall is similar to IC3. The novel aspect of our approach is the use of *diagrams* [Chang and Keisler 1990] to generalize individual states into sets of states before checking for reachability. Diagrams provide a *structural abstraction* of states by existential formulae: The *diagram* of a *finite* model σ , denoted by $\text{Diag}(\sigma)$, is an existential cube that describes explicitly the relations between all the elements of the model.⁴

Definition 4.1 (Diagrams). Given a *finite* model $\sigma = (D, \mathcal{I})$ over alphabet \mathcal{V} , the *diagram* of σ , denoted by $\text{Diag}(\sigma)$, is a closed formula over alphabet \mathcal{V} that denotes the set of models in which σ can be isomorphically embedded. $\text{Diag}(\sigma)$ is constructed as follows.

- For every element $e_i \in D$, a fresh variable x_{e_i} is introduced.
- $\varphi_{\text{distinct}}$ is a conjunction of inequalities of the form $x_{e_i} \neq x_{e_j}$ for every pair of distinct elements $e_i \neq e_j$ in the model.
- $\varphi_{\text{constants}}$ is a conjunction of equalities of the form $c = x_e$ for every constant symbol c such that $\sigma \models c = e$.
- φ_{atomic} is a conjunction of atomic formulae that include for every predicate $p \in \mathcal{V}$ the atomic formula $p(\bar{x}_e)$ if $\sigma \models p(\bar{e})$ and $\neg p(\bar{x}_e)$ otherwise.

Thus: $\text{Diag}(\sigma) \stackrel{\text{def}}{=} \exists x_{e_1} \dots x_{e_{|D|}}. \varphi_{\text{distinct}} \wedge \varphi_{\text{constants}} \wedge \varphi_{\text{atomic}}.$

Intuitively, one can think of $\text{Diag}(\sigma)$ as the formula produced by treating individuals in σ as existentially quantified variables and explicitly encoding the interpretation of every constant and every predicate using a conjunction of equalities, inequalities, and

⁴Definition 4.1, as well as the property formulated by Lemma 4.5, are an adaptation of the standard model-theoretic notion of a diagram [Chang and Keisler 1990].




Fig. 3. Graphical depiction of the model σ_b and models $\sigma_b^1, \sigma_b^2, \sigma_b^3$ of $\text{Diag}(\sigma_b)$.

atomic formulae. Note that $\text{Diag}(\sigma)$ is well defined, because we consider only models with a finite domain.

Example 4.2. The diagram of σ_b , depicted in Figure 3(σ_b), is

$$\begin{aligned}
 \text{Diag}(\sigma_b) \stackrel{\text{def}}{=} & \exists x_0, x_1, x_2. x_0 \neq x_1 \wedge x_0 \neq x_2 \wedge x_1 \neq x_2 \wedge \\
 & h = x_0 \wedge j = x_1 \wedge i = x_2 \wedge \text{null} = x_2 \wedge \\
 & \neg \text{ok}(x_0) \wedge \neg \text{ok}(x_1) \wedge \neg \text{ok}(x_2) \wedge \\
 & n^*(x_0, x_0) \wedge n^*(x_1, x_1) \wedge n^*(x_2, x_2) \wedge n^*(x_0, x_1) \wedge \\
 & \neg n^*(x_0, x_2) \wedge \neg n^*(x_1, x_0) \wedge \neg n^*(x_1, x_2) \wedge \neg n^*(x_2, x_0) \wedge \neg n^*(x_2, x_1).
 \end{aligned}$$

The first line records the fact that the domain of σ_b consists of three elements. The second line characterizes the interpretations of all the constant symbols in σ_b . The other lines capture precisely the interpretation of predicates ok and n^* in σ_b .

We say that $\sigma_1 = (D_1, \mathcal{I}_1)$ is a *substructure* of $\sigma_2 = (D_2, \mathcal{I}_2)$ if $D_1 \subseteq D_2$ and for every $v \in \mathcal{V}$, $\mathcal{I}_1(v)$ is the restriction of $\mathcal{I}_2(v)$ to D_1 . The following lemma is well known:

LEMMA 4.3. $\sigma' \models \text{Diag}(\sigma)$ iff σ is isomorphic to a substructure of σ' .

That is, the diagram of σ abstracts away the exact number of elements in the domain of σ and, as such, provides a natural abstraction of states.

Example 4.4. In Figure 3, several models of $\text{Diag}(\sigma_b)$ are depicted. For clarity, the edges drawn correspond to n -links extracted from n^* of each structure using φ_n (recall from Table II). Note that all of them contain σ_b as a substructure. For example, in σ_b^1 , there is an additional element (v_3) representing a node pointing to the head of the list, as well as an additional element (v_4) representing an additional list with a single element. In σ_b^2 , there is an additional element (v_3) representing an additional node in the list between h and j (represented by elements v_0 and v_1 , respectively). To see why σ_b^2 contains σ_b as a substructure, recall that the vocabulary contains n^* , and not n itself; while no n -annotated edge appears in σ_b^2 from v_0 to v_1 , $n^*(v_0, v_1)$ does hold in σ_b^2 as well due to the transitive nature of n^* . Similarly, σ_b^3 , which represents a list that contains two additional nodes between h and j , also contains σ_b as a substructure. As these examples demonstrate, for linked list programs modelled with n^* , the diagram-based abstraction allows us to “forget” the exact length of list segments.

The following property of diagrams will be useful in the sequel.

LEMMA 4.5. Let σ be a model over \mathcal{V} , and let φ be a closed existential first-order formula over \mathcal{V} . If $\sigma \models \varphi$, then $\text{Diag}(\sigma) \Rightarrow \varphi$.

Stated differently, $\text{Diag}(\sigma)$ is the strongest existentially quantified formula that has σ as a model. Semantically, Lemma 4.5 means that for any models σ and σ' such that $\sigma' \models \text{Diag}(\sigma)$ if $\sigma \models \varphi$, then $\sigma' \models \varphi$. This implies that if a bad state is reachable from σ

and the program can be proven correct using an inductive universal invariant \mathcal{I} , then *all* the states in σ 's diagram are unreachable, too: \mathcal{I} is an inductive invariant, and thus any state σ leading to a bad state must satisfy the (closed existential) formula $\neg\mathcal{I}$. Hence, $\text{Diag}(\sigma) \Rightarrow \neg\mathcal{I}$, which means that all states satisfying $\text{Diag}(\sigma)$ are unreachable. In this sense, the abstraction based on diagrams is precise for programs with universal invariants. This property of the diagrams will allow PDR^\forall to also prove absence of universal invariants (see Proposition 5.6).

4.2. Data Structures and Frames

PDR^\forall is shown in Algorithm 1. It uses procedures *block()* and *analyzeCEX()*, shown in Algorithm 2 and Algorithm 3, respectively, as subroutines. The algorithm uses an array F of *frames*, where a frame is a conjunction of closed universal clauses. For clarity, we refer to the i th entry of the array using subscript notation, that is, F_i instead of $F[i]$. Intuitively, frame F_i over-approximates \mathcal{R}_i , the set of i -reachable states. The algorithm also maintains a *frame counter* N that records the number of frames it developed. We refer to F_0 as the *initial* frame, to F_N as the *frontier* frame, and to any F_i , where $0 \leq i < N$, as a *back* frame.

PDR^\forall maintains several invariants that ensure that every frame F_i is an over-approximation of \mathcal{R}_i and hence that the sequence of developed frames is an over-approximation of all the traces of the program of length $N + 1$ or less. Technically, this means that the algorithm constructs an *approximate reachability sequence*.

Definition 4.6. Let $TS = (\text{Init}, \rho)$ be a transition system and \mathcal{P} a safety property. A sequence $\langle F_0, F_1, \dots, F_N \rangle$ of closed formulae is an *approximate reachability sequence* for TS and \mathcal{P} if:

- (i) $\text{Init} \Rightarrow F_0$.
- (ii) $F_i \Rightarrow F_{i+1}$, for all $0 \leq i < N$, that is, for every state σ , if $\sigma \models F_i$ then $\sigma \models F_{i+1}$.
- (iii) $F_i \wedge \rho \Rightarrow (F_{i+1})'$, for all $0 \leq i < N$, that is, for every transition $(\sigma_1, \sigma_2) \models \rho$, if $\sigma_1 \models F_i$ then $\sigma_2 \models F_{i+1}$.
- (iv) $F_i \Rightarrow \mathcal{P}$, for all $0 \leq i \leq N$.

Items (ii) and (iii) ensure that every frame includes the states of the previous frame and their successors, respectively. Together with item (i), it follows by induction that for every $0 < i \leq N$ the set of states (models) that satisfy F_i is a superset of the set \mathcal{R}_i . Furthermore, by item (iv) no frame includes a bad state.

4.3. Iterative Construction of an Approximate Reachability Sequence

PDR^\forall is an iterative algorithm. At every iteration, the algorithm either strengthens the N th frame, if it contains a bad state, or otherwise starts to develop the $N + 1$ st frame. In addition, in every iteration, it might also strengthen some of the back frames. Each strengthening of frame F_i is performed by determining a universal clause φ_i that holds for any i -reachable state, and then conjoining F_i with φ_i .

Initialization. The algorithm first checks that the initial states and the bad states do not intersect. If so, then it exits and returns the state that satisfies both *Init* and *Bad* as a counterexample (line 2). Otherwise, it sets F_0 to represent the set of initial states (line 3), F_1 to represent all possible states (line 4), and the frame counter to 1. Note that at this point, F_1 is a trivial over-approximation of the set of initial states and their successors, but it might contain bad states.

Iterative Construction. The algorithm then starts its iterative search for an inductive invariant (line 6). Recall that when the algorithm develops the N th frame, it has already managed to determine an approximate reachability sequence $\langle F_0, \dots, F_{N-1} \rangle$.

ALGORITHM 1: $\text{PDR}^\forall(\text{Init}, \rho, \text{Bad})$

```

1 if  $\text{SAT}(\text{Init} \wedge \text{Bad})$  then
2   exit invalid:  $\text{model}(\text{Init} \wedge \text{Bad})$ 
3  $F_0 := \text{Init}$ 
4  $F_1 := \text{true}$ 
5  $N := 1$ 
6 while true do
7   if there exists  $0 \leq j < N$ 
8     such that  $F_{j+1} \Rightarrow F_j$  then
9     return valid
10  if  $\neg \text{SAT}(F_N \wedge \text{Bad})$  then
11     $F_{N+1} := \text{true}$ 
12     $N := N + 1$ 
13  else
14     $\sigma_b := \text{model}(F_N \wedge \text{Bad})$ 
15     $\text{block}(N, \sigma_b)$ 

```

ALGORITHM 2: $\text{block}(j, \sigma)$

```

21  $\varphi = \text{Diag}(\sigma)$ 
22 if  $(j = 0) \vee (j = 1 \wedge \text{SAT}(\varphi \wedge \text{Init}))$  then
23    $\text{analyzeCEX}(j, N)$ 
24 while  $\text{SAT}(F_{j-1} \wedge \rho \wedge (\varphi)')$  do
25    $\sigma_a = \text{reduct}_\forall(\text{model}(F_{j-1} \wedge \rho \wedge (\varphi)'))$ 
26    $\text{block}(j - 1, \sigma_a)$ 
27 for  $i = 0 \dots j$  do
28    $F_i := F_i \wedge \neg \varphi$ 

```

ALGORITHM 3: $\text{analyzeCEX}(j, N)$

```

31 if  $j = 0 \wedge$  there exists  $\sigma_0, \dots, \sigma_N$ 
32   such that
33      $\sigma_0 \models \text{Init}$ 
34      $(\sigma_i, \sigma_{i+1}) \models \rho$  for every  $0 \leq i < N$ , and
35      $\sigma_N \models \text{Bad}$ 
36   then exit invalid:  $\sigma_0, \dots, \sigma_N$ 
37 else exit no universal invariant exists

```

Hence, every iteration starts by checking whether a fixpoint has been reached (line 7). If true, then an inductive invariant proving unreachability of *Bad* has been found, and the algorithm returns valid (line 8). Otherwise, the algorithm keeps on strengthening the frontier frame F_N by searching for a *bad witness*, a bad state in the frontier frame (line 9). If no such state exists, then it means that no bad state is N -reachable. Moreover, at this point $\langle F_0, \dots, F_N \rangle$ is an approximate reachability sequence. Thus, the iterative strengthening of F_N terminates and a new frontier frame is initialized to *true* (line 10 and 11).

If the frontier frame contains a bad witness, that is $F_N \wedge \text{Bad}$ is satisfiable, then there *might* be an N -reachable bad state. Due to our requirement for finite satisfiability of the logic, the bad witness is a *finite* model. Given a bad witness σ_b (line 13), the algorithm tries to determine whether it is indeed reachable, and thus the program does not satisfy its specification or whether σ_b was discovered due to some over-approximation in one of the back frames. This check is done by invoking procedure $\text{block}()$ with the index of the frontier frame and σ_b as parameters (line 14). The latter either returns a counterexample, determines that it is impossible to prove the specification using a universal invariant (in the given logic and vocabulary), or strengthens the frontier frame to exclude the *set of states in the diagram of* σ_b , and possibly strengthens some back frames, too (see below). The iterative construction and strengthening of the frames continues until reaching a fixpoint, finding a counterexample, or determining the absence of a universal invariant.⁵

Example 4.7. When analyzing the running example, our algorithm discovers that state σ_b , shown in Figure 4 (as well as in Figure 3), is a bad witness when $F_1 = \text{true}$, and thus it invokes $\text{block}(1, \sigma_b)$. In this example, $\text{block}()$ succeeds to block σ_b . Unfortunately, the strengthened frame F_1^1 (see below) still has bad models. Therefore, the iterative strengthening continues and the next iterations find σ'_b , depicted in Figure 4, as a bad

⁵For efficiency, in our implementation we represent each frame as a set of clauses (with the meaning of conjunction) and check implication (line 7) by checking inclusion of these sets. To facilitate this fixpoint computation, any clause φ in F_i that is *inductive* in F_i , that is, $F_i \wedge \rho \Rightarrow (\varphi)'$ is also propagated forward to F_{i+1} . In particular, this allows to initialize a new frontier frame F_N , for $1 < N$, to a tighter over-approximation of \mathcal{R}_N than *true* (line 10) [Bradley 2011].




Fig. 4. Graphical depiction of models found during the analysis of the running example.

witness model for F_1^1 , σ''_b as a bad witness model of F_1^2 and σ'''_b as a bad witness model of F_1^3 . At that point, however, the algorithm determines that the strengthened frame F_1^4 does not have a bad witness. $\langle F_0, F_1^4 \rangle$ is now an approximate reachability sequence, and PDR^\forall goes on and initializes a new frame, F_2 , to *true*, and the search for an inductive invariant continues.

Diagram-Based Abstract Blocking. Procedure $\text{block}(j, \sigma)$, shown in Algorithm 2, gets an index of a frame $j = 0, \dots, N$ and a state σ that is included in the j th frame, that is, $\sigma \models F_j$, and tries to determine whether σ is j -reachable. The unique aspect of our approach is the way in which it abstracts σ to a set of states in order to accelerate the strengthening routine. Namely, the use of diagrams. More specifically, PDR^\forall computes the diagram φ of σ (line 21) and then checks whether there is a j -reachable state satisfying φ . Importantly, due to Lemma 4.5, if a universal invariant exists, then the generalization of σ to its diagram will not include any reachable state, and hence the abstraction is precise in the sense that it maintains unreachability. In this case, the strengthening of F_j is also guaranteed to succeed, excluding not only σ , but its entire diagram.

The check if the diagram φ of σ includes a j -reachable state is done *conservatively* by determining whether some state of φ is an initial state or has a predecessor in F_{j-1} . (Recall that F_{j-1} over-approximates \mathcal{R}_{j-1} .) The former is equivalent to checking if $\varphi \wedge \text{Init}$ is satisfiable. Note that if we reached the initial frame, that is, if $j = 0$, then $\sigma \models \text{Init}$, and hence the above formula is guaranteed to be satisfiable. Explicitly checking that $\varphi \wedge \text{Init}$ is satisfiable is required only at the second frame, that is, if $j = 1$ (see proof in Section 5.1):

LEMMA 4.8. *For every $1 < j \leq N$, when $\text{block}(j, \sigma)$ is called, $F_i \Rightarrow \neg \text{Diag}(\sigma)$ for every $i \leq j - 1$. In particular, $\text{Init} \Rightarrow \neg \text{Diag}(\sigma)$.*

If the algorithm finds an *adverse initial state*, that is, an initial state satisfying φ , (line 22),⁶ then it invokes procedure $\text{analyzeCEX}()$ for further analysis (see below). Otherwise, the algorithm checks if the formula $\delta = F_{j-1} \wedge \rho \wedge (\varphi)'$ is satisfiable (line 24),⁷ that is, whether some state of φ has a predecessor in F_{j-1} . There can be two cases:

Case I. If δ is unsatisfiable, then no state represented by φ is j -reachable. Hence, F_j remains an over-approximation of \mathcal{R}_j even if any state of φ is excluded. The exclusion

⁶If Init is a universal formula, then Lemma 4.8 holds for $j = 1$ as well, and, hence, $j = 1 \wedge \text{SAT}(\varphi \wedge \text{Init})$ never holds, and its check can be omitted (line 22).

⁷As an optimization, one can consider $\delta' = F_{j-1} \wedge \neg \varphi \wedge \rho \wedge (\varphi)'$ instead of δ . The two formulae are equivalent, since $F_{j-1} \Rightarrow \neg \varphi$ (by Lemma 4.8 for $j > 1$, and since it was checked for $j = 1$), but the strengthening of δ can make the satisfiability check cheaper.

is done by conjoining the j th frame with the universal formula $\neg\varphi$ (line 28) and results in a strengthening of F_j . In fact, $\neg\varphi$ is conjoined to any back frame (line 27). Similarly to traditional PDR, the last step is done in order to maintain the inclusion property of frames (Definition 4.6(ii)). Moreover, this update maintains all the properties of Definition 4.6, in particular, $F_i \wedge \rho \Rightarrow (F_{i+1})'$ (iii) is preserved despite conjoining $(F_{i+1})'$ with $(\neg\varphi)'$, because frame inclusion ensures that for every $i < j$, $F_i \wedge \rho \Rightarrow F_{j-1} \wedge \rho$ and the latter implies $(\neg\varphi)'$. Note that preserving the properties of the approximate reachability sequence means that after the update, F_i remains an over-approximation of \mathcal{R}_i . In the following, we refer to the exclusion of the states of φ as the *blocking* of (the diagram of) σ from frame F_j .

Example 4.9. In our running example, in the first iteration $\text{block}(1, \sigma_b)$ updates F_1^0 to $F_1^1 = \text{true} \wedge \neg\text{Diag}(\sigma_b)$. This excludes σ_b , but also all states where $i = \text{null}$, ok is empty, and j is n -reachable from h in *any* (nonzero) number of steps (e.g., the states σ_b^2 and σ_b^3 depicted in Figure 3). In later iterations block updates $F_1^2 = F_1^1 \wedge \neg\text{Diag}(\sigma_b')$, $F_1^3 = F_1^2 \wedge \neg\text{Diag}(\sigma_b'')$, and $F_1^4 = F_1^3 \wedge \neg\text{Diag}(\sigma_b''')$.

Case II. If δ is satisfiable, then there exists an *adverse* state σ_a in frame F_{j-1} , a state that is the predecessor of some state of the diagram of σ that we try to block at frame F_j . Note that σ_a is not necessarily a predecessor of σ itself. The adverse state σ_a is found by taking the reduct of a (finite) model of δ to \mathcal{V} (line 25). If an adverse model σ_a exists, then the algorithm *recursively* tries to block it from F_{j-1} (line 26). The recursive procedure continues until the adverse state is either blocked or the algorithm finds an adverse initial state (line 22). Note that blocking an adverse state during the development of the N th frame leads to a strengthening of some back frame F_i and thus tightens its over-approximation of \mathcal{R}_i .

Finding Concrete Counterexamples and Proving the Absence of Universal Invariants. Procedure $\text{analyzeCEX}()$, shown in Algorithm 3, is called when an adverse initial state is found. Such a state indicates that an abstract counterexample exists:

Definition 4.10 (Abstract and Spurious Counterexamples). A sequence of closed formulae $\langle \varphi_j, \varphi_{j+1}, \dots, \varphi_N \rangle$ is an *abstract counterexample* if the formulae $\varphi_j \wedge \text{Init}$, $\varphi_N \wedge \text{Bad}$, and $\varphi_i \wedge \rho \wedge (\varphi_{i+1})'$, for every $i = j, \dots, N-1$, are all satisfiable. The abstract counterexample is *spurious* if there exists no sequence of states $\langle \sigma_j, \sigma_{j+1}, \dots, \sigma_N \rangle$ such that $\sigma_j \models \text{Init}$, $\sigma_N \models \text{Bad}$, and for every $j \leq i < N$, $(\sigma_i, \sigma_{i+1}) \models \rho$.

An abstract counterexample does not necessarily describe a real counterexample. In fact, if $j \neq 0$, the counterexample is necessarily spurious (as, if a real counterexample shorter than N had existed, the algorithm would have already terminated during the development of the $N-1$ th frame). However, when $j = 0$, the algorithm determines if the abstract counterexample is real or spurious by checking whether a bad state can be reached by N applications of the transition relation (line 31). Technically, $\text{analyzeCEX}()$ can be implemented using a symbolic bounded model checker [Biere et al. 2003]. If a real counterexample is found, then the algorithm reports it (line 35). Otherwise, the obtained counterexample is *spurious*. Technically, this means that the property is neither verified nor falsified. In our case, the algorithm can determine that the verification effort is doomed: The spurious counterexample is in fact a proof for the absence of a universal invariant (see Prop. 5.6).

Generalization of Blocked Diagrams. Rather than blocking a diagram φ from frames $0, \dots, j$ by conjoining them with the clause $\neg\varphi$ (line 28), our implementation uses a minimal unsat core of $\psi = ((\text{Init})' \vee (F_{j-1} \wedge \rho)) \wedge (\varphi)'$ to define a clause L that implies $\neg\varphi$

and is also disjoint from $Init$ and unreachable from F_{j-1} . Blocking is done by conjoining L with F_i for every $i \leq j$.⁸

5. CORRECTNESS

In this section, we formalize the correctness guarantees of PDR^\forall . We start by formalizing the invariants of the sequence of frames maintained by PDR^\forall . We then prove that the output of PDR^\forall is correct.

5.1. Properties of the Frames Computed by PDR^\forall

The following lemma summarizes the invariants of the sequence of frames computed by PDR^\forall . It follows from a simple induction on the steps of PDR^\forall . These invariants are adopted from traditional PDR .

LEMMA 5.1. *Let $TS = (Init, \rho)$ be a transition system and \mathcal{P} a safety property. For every $N \geq 1$, the sequence $\langle F_0, F_1, \dots, F_N \rangle$ obtained by PDR^\forall right before N is increased in Line 11 is an approximate reachability sequence for TS and \mathcal{P} (Definition 4.6). Further, in every other step of the while loop (Line 6), the sequence $\langle F_0, F_1, \dots, F_N \rangle$ satisfies all the requirements of Definition 4.6, except for, possibly, requirement (iv) for $i = N$.*

In order to provide a slightly tighter characterization of the frames computed by PDR^\forall , which, unlike Lemma 5.1, is specific to PDR^\forall , we need the following definition.

Definition 5.2 (Relaxed Traces). Given a transition system $TS = (Init, \rho)$, we say that a sequence of states $\sigma_0, \dots, \sigma_n$ is a *relaxed trace* of TS if for every $0 \leq i \leq n-1$, there exists σ'_{i+1} such that $(\sigma_i, \sigma'_{i+1}) \models \rho$ and $\sigma'_{i+1} \models (Diag(\sigma_{i+1}))'$.

For a transition system TS and a property \mathcal{P} , with $Bad = \neg\mathcal{P}$, we denote by B_i the set of states that can reach a bad state in at most i steps via a relaxed trace. That is, $\sigma \in B_i$ if and only if there exists a relaxed trace $\sigma_0, \dots, \sigma_n$ such that $\sigma_0 = \sigma$, $\sigma_n \models Bad$, and $n \leq i$. In particular, $B_0 = \{\sigma \mid \sigma \models Bad\}$. We denote by B the set of states that can reach a bad state in some number of steps via a relaxed trace, that is, $B = \bigcup_{i \geq 0} B_i$.

The following lemma states that the frames computed by PDR^\forall do not contain states that lead to bad states via relaxed traces, where the length of the considered traces depends on the frame index.

LEMMA 5.3. *Let $TS = (Init, \rho)$ be a transition system and \mathcal{P} a safety property. For every $N \geq 1$ and $0 \leq j \leq N-1$, if $\sigma \models F_j$, then $\sigma \notin B_{N-1-j}$.*

That is, F_j does not include any state that reaches a bad state via a relaxed trace in $N-1-j$ steps or fewer. Note that this lemma holds in particular when N is increased to $N+1$ (Line 11), in which case it implies that for every $0 \leq j \leq N$, if $\sigma \models F_j$, then $\sigma \notin B_{N-j}$. The proof follows a simple induction on j that shows that if $\sigma \in B_{N-1-j}$, then it is excluded from F_j .

On the other hand, it is easy to prove by induction that every state that PDR^\forall attempts to block in frame F_j is a state in B .

LEMMA 5.4. *Let $TS = (Init, \rho)$ be a transition system and \mathcal{P} a safety property. For every $N \geq 1$ and $0 \leq j \leq N$, if $block(j, \sigma)$ is called, then $\sigma \in B_{N-j}$.*

We can now return to the proof of Lemma 4.8.

⁸We can also use inductive generalization, that is, look for a minimal subclause L of $\neg\varphi$ that is still inductive relative to F_{j-1} , meaning $((Init)' \vee (F_{j-1} \wedge L \wedge \rho)) \wedge (\neg L)'$ is unsatisfiable.

PROOF OF LEMMA 4.8. Suppose $\sigma \models F_j$ is discovered in the backward traversal from F_N . By Lemma 5.4, $\sigma \in B_{N-j}$. Therefore, by Lemma 5.3, $\sigma \not\models F_{j-1}$. To complete the proof, recall that we consider $j > 1$. Therefore $j - 1 > 0$, hence F_{j-1} is a universal formula. Therefore, by Lemma 4.5, $\text{Diag}(\sigma) \Rightarrow \neg F_{j-1}$, or, equivalently, $F_{j-1} \Rightarrow \neg \text{Diag}(\sigma)$. Since $F_i \Rightarrow F_{j-1}$ for every $i \leq j - 1$, we conclude that $F_i \Rightarrow \neg \text{Diag}(\sigma)$ for every $i \leq j - 1$. \square

Note that if *Init* is a universal formula, then the claim formulated by Lemma 4.8 holds for $j = 1$ as well. That is, when $\text{block}(1, \sigma)$ is called, $F_0 \Rightarrow \neg \text{Diag}(\sigma)$. In this case, the additional check of $(j = 1 \wedge \text{SAT}(\varphi \wedge \text{Init}))$ performed in Line 22 of $\text{block}(j, \sigma)$ is always false and hence not needed.

5.2. Correctness of the Outcome of PDR^\forall

We recall that if PDR^\forall terminates it reports that either the program is safe, the program is not safe, providing a counterexample, or the program cannot be verified using a universal inductive invariant.

LEMMA 5.5. *Let $TS = (\text{Init}, \rho)$ be a transition system and let \mathcal{P} be a safety property. If PDR^\forall returns valid, then TS satisfies \mathcal{P} . Further, if PDR^\forall returns a counterexample, then TS does not satisfy \mathcal{P} .*

PROOF. PDR^\forall returns valid if there exists i such that $F_{i+1} \Rightarrow F_i$. Therefore, $F_i \wedge \rho \Rightarrow (F_{i+1})' \Rightarrow (F_i)'$. Recall that, by the properties of an approximate reachability sequence, $\text{Init} \Rightarrow F_0 \Rightarrow F_i$ and $F_i \Rightarrow \mathcal{P}$. Therefore, F_i is an inductive invariant, which ensures that TS satisfies \mathcal{P} . The second part of the claim follows immediately from the definition of a counterexample. \square

Note that the proof of Lemma 5.5 also implies that when PDR^\forall determines that TS satisfies \mathcal{P} , it also obtains a *universal* inductive invariant. Namely, the frame F_i in which a fixpoint was identified comprises such an invariant. (Recall that each frame is a universal formula, as it is obtained as a conjunction of negated diagrams.)

PROPOSITION 5.6. *Let $TS = (\text{Init}, \rho)$ be a transition system and let \mathcal{P} be a safety property. If PDR^\forall obtains a spurious counterexample $\langle \varphi_j, \dots, \varphi_N \rangle$, then there exists no universal safety inductive invariant \mathcal{I} for TS and \mathcal{P} over the given vocabulary.*

In order to prove Proposition 5.6, we first show that if a universal inductive invariant exists for TS and \mathcal{P} , then no state that reaches a bad state via a relaxed trace satisfies the invariant:

PROPOSITION 5.7. *Let $TS = (\text{Init}, \rho)$ be a transition system and \mathcal{P} a safety property, and let B be the set of states that reach a bad state via a relaxed trace (see Section 5.1). Let \mathcal{I} be a universal inductive safety invariant for \mathcal{P} . For any $\sigma \in B$, we have $\sigma \models \neg \mathcal{I}$.*

PROOF. We show by induction on i that $\sigma \in B_i$ implies $\sigma \models \neg \mathcal{I}$. For the base case, $\sigma \in B_0$, we have $\sigma \models \text{Bad}$. Since $\text{Bad} = \neg \mathcal{P}$ and $\mathcal{I} \Rightarrow \mathcal{P}$, we conclude that $\sigma \models \neg \mathcal{I}$.

For the inductive step, let $\sigma \in B_{i+1}$. By definition of B_{i+1} , there exist models σ_i and σ'_i such that $(\sigma, \sigma'_i) \models \rho \wedge (\text{Diag}(\sigma_i))'$ and $\sigma_i \in B_i$. Moreover, by the induction hypothesis, $\sigma_i \models \neg \mathcal{I}$. Since $\neg \mathcal{I}$ is an existential formula, this means by Lemma 4.5 that $\text{Diag}(\sigma_i) \Rightarrow \neg \mathcal{I}$. We conclude that $\rho \wedge (\text{Diag}(\sigma_{i+1}))' \Rightarrow \rho \wedge (\neg \mathcal{I})'$. Therefore, (σ, σ'_i) is also a model of the formula $\rho \wedge (\neg \mathcal{I})'$.


If we assume that $\sigma \models \mathcal{I}$, then we would get that $\mathcal{I} \wedge \rho \wedge (\neg \mathcal{I})'$ is satisfiable, in contradiction \mathcal{I} being inductive. Hence, $\sigma \models \neg \mathcal{I}$. \square


```

requires:  $p = \text{null} \wedge q = \text{null} \wedge i = g \wedge g \neq \text{null} \wedge j = h \wedge h \neq \text{null} \wedge$ 
 $\exists v. n^*(g, v) \wedge n^*(h, v) \wedge v \neq \text{null}$ 
ensures:  $p = q \wedge p \neq \text{null} \wedge i = \text{null} \wedge j = \text{null}$ 
void traverseTwo(List g, List h) {
  i := g; j := h;
  while (i  $\neq$  null  $\vee$  j  $\neq$  null) {
    if i  $\neq$  null then { p := i; i := i.n };
    if j  $\neq$  null then { q := j; j := j.n };
  }
}

```

Fig. 5. A procedure that finds the last elements of two non-empty acyclic lists.

Fig. 6. A spurious counterexample found for procedure `traverseTwo()`, shown in Figure 5.

PROOF OF PROPOSITION 5.6 Assume that there exists a universal safety inductive invariant \mathcal{I} over \mathcal{V} . By Lemma 5.4, for every state σ_i generated by PDR^\forall at frame F_i , $\sigma_i \in B$. Hence by Proposition 5.7, $\sigma_i \models \neg \mathcal{I}$. This implies, by Lemma 4.5, that every diagram φ_i generated by PDR^\forall at frame F_i is such that $\varphi_i \Rightarrow \neg \mathcal{I}$, and hence $\varphi_i \Rightarrow \neg \text{Init}$. (Recall that by definition $\text{Init} \Rightarrow \mathcal{I}$, that is, $\neg \mathcal{I} \Rightarrow \neg \text{Init}$). This contradicts the existence of a spurious counterexample, where $\varphi_j \wedge \text{Init}$ is satisfiable. \square

Example 5.8. Procedure `traverseTwo()`, presented in Figure 5 together with its pre- and post-condition, traverses two lists until it finds their last elements. If the lists have a shared tail, then p and q should point to the same element when the traversal terminates. The program indeed satisfies this property. However, this cannot be proven correct using an inductive universal invariant: Take, as usual, Init to be the procedure's precondition and \mathcal{P} to be the safety property whose negation is $\text{Bad} = (i = \text{null} \wedge j = \text{null}) \wedge \neg \text{Post}$, where Post is the procedure's postcondition. Consider the state σ_0 depicted in Figure 6. Clearly, this model satisfies Init . Therefore, if \mathcal{I} exists, $\sigma_0 \models \mathcal{I}$. σ_0 is a predecessor of σ_1^t , and hence it should be the case that $\sigma_1^t \models \mathcal{I}$. Now consider σ_1 , which is a submodel of σ_1^t and interprets all constants as in σ_1 . If \mathcal{I} is universal, then $\sigma_1 \models \mathcal{I}$ as well. The model σ_2 is a successor of σ_1 and hence $\sigma_2 \models \mathcal{I}$. However, $\sigma_2 \not\models \mathcal{P}$, in contradiction to the property of a safety invariant. Indeed, when using PDR^\forall , the spurious counterexample $\langle \sigma_0, \sigma_1, \sigma_2 \rangle$ presented in Figure 6 is obtained. This indicates that no universal invariant for \mathcal{P} exists. Note that state σ_1 is a predecessor of σ_2 and recall that σ_0 is a predecessor of σ_1^t . The spurious counterexample was obtained, because σ_1^t satisfies the diagram of state σ_1 .

6. SUFFICIENT CONDITIONS FOR TERMINATION

The PDR^\forall algorithm is in general not guaranteed to terminate: It does not restrict the number of distinct existentially quantified variables in diagrams. This comes as a blessing for finding invariants with an arbitrary number of quantified variables but is a curse when it comes to bounding the search space. There are, however, non-trivial classes of programs where it terminates. In particular, we show that PDR^\forall terminates

for programs that manipulate singly linked lists. Inferring universal invariants for this class of programs was shown to be decidable [Padon et al. 2016, Theorem 4.2] using the naïve backwards reachability algorithm presented in Algorithm 4. The main result in this section bridges backwards reachability with PDR^\forall , here formulated for a transition system $TS = (\text{Init}, \rho)$ and a property \mathcal{P} , with $\text{Bad} = \neg\mathcal{P}$:

PROPOSITION 6.1. *If backwards reachability, given by Algorithm 4, terminates on input transition system TS and property \mathcal{P} , then PDR^\forall terminates as well.*

In order to establish this result, we introduce a notion of *effective* encodings of the (backwards) reachable states, B (defined in Section 5.1). In a nutshell, a set of (backwards) reachable states B is effective for a class of transition systems and properties if the set of reachable states can be defined using a finitary existential formula, that is, a formula that uses a finite number of quantified variables and contains a finite number of sub-formulae. This allows us to connect backwards reachability with PDR^\forall using the two lemmas:

LEMMA 6.2. *If Algorithm 4 terminates on TS and \mathcal{P} , then B is effective.*

and

LEMMA 6.3. *If B is effective for TS and \mathcal{P} , then PDR^\forall is guaranteed to terminate.*

We start with the description of Algorithm 4. It computes the complement of the least fixpoint of diagrams that can reach (via a relaxed trace) the bad states Bad . If the states represented by the resulting formula do not properly contain the initial states Init , then there are no universal inductive invariants, otherwise there is one.

ALGORITHM 4: Naïve Backward Reachability

```

1  $I := \text{true}$ 
2 while  $\text{SAT}(I \wedge \text{Bad})$  or  $\text{SAT}(I \wedge \rho \wedge (\neg I))$  do
3    $(\sigma, \sigma') := \text{model}((I \wedge \text{Bad}) \vee (I \wedge \rho \wedge (\neg I)))$ 
4    $I := I \wedge \neg \text{Diag}(\sigma)$ 
5 if  $\text{SAT}(\text{Init} \wedge \neg I)$  then
6   return no universal inductive invariant
7 else
8   return  $I$  is a universal inductive invariant

```

Next, let us define what we mean by a set of backward reachable states being *effective*. Recall that we use B_i to denote the set of states that can reach a bad state via a relaxed trace of TS in at most i steps and B to denote the set of states that can reach a bad state via a relaxed trace in some number of steps. Thus, $B = \bigcup_{i \geq 0} B_i$ (see Section 5.1).

Definition 6.4. We say that B is *effective* if it can be described by an existential formula. That is, there exists an existential formula ψ_B such that $B = \{\sigma \mid \sigma \models \psi_B\}$.

Lemma 6.2 follows from the description of the naïve algorithm and the definition of effective states:

PROOF OF LEMMA 6.2. Suppose Algorithm 4 terminates, and let $\neg I$ be the negation of the obtained formula, I . Then $B = \{\sigma \mid \sigma \models \neg I\}$. Further, since $\neg I$ is equivalent to a disjunction of diagrams, it is an existential formula. This ensures effectiveness. \square

Establishing Lemma 6.3 requires a bit more context. Recall that we require that every satisfiable formula in \mathcal{L} has a finite model and assume to have a decision procedure

$SAT(\psi)$, which checks if a formula ψ in \mathcal{L} is satisfiable, and a function $model(\psi)$, which returns a *finite* model σ of ψ if such a model exists and `None` otherwise. For the termination argument, we place a stronger requirement: that finite model sizes are functions of the number of constants and existentially quantified variables. In other words, we require the existence of a function $bound : \mathbb{N} \rightarrow \mathbb{N}$ that given a formula ψ in \mathcal{L} , with at most n_ψ existentially quantified variables and constants, ψ has a finite model if and only if it has a model that contains no more than $bound(n_\psi)$ elements. Note that such a bound only depends on the existential quantifiers and constants, and not, for example, on the number of universal quantifiers. We note that *EPR* has such a bound; see Section 6.1.

Further, we also assume that $model(\psi)$ always returns a model of size at most $bound(n_\psi)$. Assuming that \mathcal{L} satisfies the additional requirement, and that all satisfiability checks performed by PDR^\forall are in \mathcal{L} , we can prove Lemma 6.3.

PROOF OF LEMMA 6.3. The proof consists of two main arguments. First, we show that there exists a window of frames of a fixed length in which PDR^\forall generates new clauses. Then, we show that the size of generated clauses is bounded by a function that depends on their distance from the last frame. As the latter is bounded by the length of the window, we obtain a bound on the size of clauses, which ensures termination.

We first note that, since B is effective, there exists k such that $B = B_k$. This holds, because every existential formula can be written as a finite disjunction of diagrams. Indeed, this can be achieved by converting the formula to DNF and performing for every existential cube a case splitting on whether the existentially quantified variables are distinct or not and on whether the missing instances of constants and relations appear negatively or positively. Therefore, effectiveness of B implies that there exists a finite set of states $\sigma_1, \dots, \sigma_m$ such that $B = \{\sigma \mid \sigma \models \bigvee_{i=1}^m Diag(\sigma_i)\}$. For each $i = 1, \dots, m$, let k_i be the length of a shortest relaxed trace leading from σ_i to a bad state, and let $k = \max_{i=1}^m k_i$. Since every state in B is a model of $Diag(\sigma_i)$ for at least one of the states σ_i , every such state starts a relaxed trace of length at most k_i to a bad state. Therefore, $B = B_k$. Hence for every $i > k$, $B_i = B_k$ as well.

We now turn to show that in each step, PDR^\forall only generates new clauses in the last k frames. Consider a fixed N . By Lemma 5.3, it holds that for every $j < N$, if $\sigma \models F_j$, then $\sigma \notin B_{N-1-j}$. By the choice of k , if $N - 1 - j \geq k$, then $B_{N-1-j} = B$. This means that for every $j \leq N - 1 - k$, if $\sigma \models F_j$, then $\sigma \notin B$. On the other hand, by Lemma 5.4, every state that PDR^\forall attempts to block in frame F_j is a state in B . As no such state exists for $j \leq N - 1 - k$, we conclude that no state is blocked at F_j for $j \leq N - 1 - k$. Therefore, new clauses, which result from blocked states, are only generated in the last $k + 1$ frames F_{N-k} to F_N . (They are, of course, pushed backwards once generated.)

We now show that we can bound the size of models obtained by PDR^\forall in its backward traversal and hence can bound the size of generated clauses. Let n_c denote the number of constants in \mathcal{V} , n_{Bad} denote the number of existential quantifiers in *Bad*, and n_ρ denote the number of existential quantifiers in ρ . We show by induction on j that we can bound the size of models obtained by PDR^\forall in frame $N - j$ by a function of n_c , n_{Bad} , and n_ρ (i.e., the bound does not depend on additional parameters such as N). For the base case ($j = 0$) recall that the backward traversal starts from a state $\sigma \models F_N \wedge Bad$. Since F_N is a universal formula, the properties of \mathcal{L} ensure that the size of the obtained model is bounded by $bound(n_c + n_{Bad})$. For the induction step, when PDR^\forall makes a step backward from a diagram of a state σ using ρ , it uses the formula $F_{j-1} \wedge \rho \wedge (\varphi)'$, where $\varphi = Diag(\sigma)$. The only existential quantifiers in this formula result from ρ and φ , where the number of the latter is equal to the size of the domain of σ (by the construction of the diagram). Our assumptions on \mathcal{L} and on $model(\psi)$ therefore ensure that the size of

the domain of the obtained model is bounded by

$$\text{bound}(n_c + n_\rho + n), \quad (*)$$

where n is the size of σ . By the induction hypothesis, n is bounded by a function of n_c , n_{Bad} , and n_ρ only, and, hence, the claim follows.

We denote by f_j the function of n_c , n_{Bad} , and n_ρ that provides a bound on the size of models obtained by PDR^\forall in frame $N - j$. Therefore, when PDR^\forall makes at most k backward steps from Bad , the number of elements in the obtained models is bounded by

$$\max = \max_{j=0}^k f_j(n_c, n_{Bad}, n_\rho).$$

This provides a bound on the number of elements in the models that PDR^\forall tries to block.

Altogether, we conclude that the clauses generated by PDR^\forall (blocked diagrams) have at most \max quantifiers, which makes the potential number of clauses finite and hence ensures termination of PDR^\forall . \square

6.1. Termination When Reasoning with Effectively Propositional Logic

If $Init$ and ρ are *EPR* formulae and \mathcal{P} is a universal formula, then all the satisfiability queries made by PDR^\forall are of *EPR* formulae. *EPR* has the finite model property, and its satisfiability is decidable. Further, *EPR* meets the additional requirement needed for termination, as an *EPR* formula ψ is satisfiable if and only if it has a satisfying model whose size is bounded by n_ψ , where n_ψ is the number of constants and existentially quantified variables in ψ . That is, for every *EPR* formula ψ , $\text{bound}(n_\psi) = n_\psi$. The existence of this bound is a well-known property of the Bernays-Schönfinkel-Ramsey class of first-order formulae [Börger et al. 2008].

Note that, in this case, the bound on the number of quantifiers in clauses generated by PDR^\forall in the last $k + 1$ frames can be explicitly expressed as $(k + 1) \cdot n_c + k \cdot n_\rho + n_{Bad}$, where n_c denotes the number of constants in \mathcal{V} , n_{Bad} denotes the number of existential quantifiers in Bad , and n_ρ denotes the number of existential quantifiers in ρ . This can be proven by induction on k using the bound (*).

In Padon et al. [2016], it is shown that Algorithm 4 is guaranteed to terminate in the case where the substructure relation is a well-quasi-order on the set of states.⁹ This is the case for programs manipulating singly linked lists that are modeled in *EPR* using a single transitive reflexive binary relation for n^* (axiomatized as in Itzhaky et al. [2013]), any number of constants and any number of unary relations (but no additional binary or higher-arity relations) [Padon et al. 2016]. We therefore conclude that PDR^\forall also terminates on such programs.

7. IMPLEMENTATION AND EMPIRICAL EVALUATION

PDR^\forall is parametric in the vocabulary and can be implemented on top of any decision procedure for finite satisfiability of first-order logic formulae. The language of these formulae should be expressive enough to capture the assertions, transition system, and space of candidate invariants. Our algorithm is not guaranteed to terminate, thus the underlying logic does not have to be decidable. Our implementation, however, uses EA^R as explained in Section 3.

⁹A *well-quasi-ordering* \sqsubseteq on a set X is a preorder (i.e., a reflexive, transitive binary relation) such that any infinite sequence of elements x_0, x_1, \dots from X contains an increasing pair $x_i \sqsubseteq x_j$ with $i < j$.

We note that the use of EA^R is key for the ability of PDR^\forall to infer inductive invariants of list-manipulating programs through generalization of particular counterexamples: At first blush, one might expect the tool to keep enumerating formulae about lists of every possible length. Luckily, EA^R specifies transition relations and properties using n^* instead of n . This formulation gives a natural abstraction for the lengths of list segments. In combination with generalization, via diagrams and unsat cores, this allows inferring clauses that apply to many lengths. In particular, we can reason about reachability without having to enumerate all the possible lengths. For example, consider Example 4.7: When blocking the diagram of the state σ_b presented in Figure 4, we also block all lists in which j is reachable from h in any number of steps.

Benchmarks. We implemented¹⁰ PDR^\forall and applied it to a collection of procedures that manipulate singly linked lists, doubly linked lists, multi-linked lists, and implementations of an insertion-sort algorithm [Cormen et al. 2009], and a union-find algorithm [Cormen et al. 2009]. Our experiments were conducted using a 3.6GHz Intel Core i7 machine with 32GB of system memory, running Ubuntu 14.04. We used the 64bit version of Z3 4.4.0 (build hashcode 0482e7fe727c) [de Moura and Bjørner 2008] with the default settings to check satisfiability of EPR formulae. Table V summarizes our experimental results.

(a) *Verification.* Our analyzer successfully verified (i) *memory safety*, that is, the absence of null-dereferences and of memory leaks; (ii) *preservation of data-structure integrity*, meaning that the procedure never creates cycles in the list (for acyclic list manipulating procedures) or that the cyclic structure is preserved (for procedures operating on cyclic lists); and (iii) *functional correctness*, for several singly and doubly linked list manipulating procedures and procedures operating with (restricted) cyclic lists. The precondition says that the expected input is a (possibly empty) (a-)cyclic list, and the post-condition is the one expected from the procedure’s name. For example, the post-condition of `reverse()` is that it returns a list comprised of the same elements as in its input but in reversed order.

To verify the absence of memory leaks, we used a unary predicate $alloc(\cdot)$ to record whether a node is allocated. We encode the instructions `new` or `free` by means of $alloc(\cdot)$ by updating it accordingly in the transition relation. The absence of memory leaks then is formulated, informally, by saying that all non-reachable elements do not satisfy $alloc$: $\forall \alpha. alloc(\alpha) \rightarrow \bigvee_{x \in PVar} n^*(x, \alpha)$, where $PVar$ denotes the set of variables in the program. Additionally, we used auxiliary (ghost) constants and predicates and to mark the elements of the input list and record the reachability order between them. For example, we can specify a reversed order condition by introducing a ghost predicate \underline{n}^* for the initial reachability and using the following formula as the post-condition: $\forall \alpha \beta. \underline{n}^*(h, \alpha) \wedge \underline{n}^*(\alpha, \beta) \leftrightarrow n^*(\beta, \alpha) \wedge n^*(\alpha, h)$.

Our current implementation uses a fault avoidance version of *wlp* that ensures that the programs do not perform null-pointer dereferences and preserve the data structure integrity. Checking for the absence of memory leaks is done by adding the aforementioned formula to the post-condition. Thus, the user is required only to come up with the post-conditions pertaining to functional correctness.

We also verified the correctness of several procedures that manipulate sorted lists: `sorted-insert()` inserts an element into its appropriate place in a sorted list, `sorted-merge()` creates a sorted list by merging two sorted ones, and `bubble-sort()` and `insertion-sort()` sort their input lists. We represented the order on data elements

¹⁰The implementation is available for download at <https://bitbucket.org/tausigplan/updr-distrib/>.

Table V. Experimental Results. Running Time Is Measured in Seconds

	Full				Memory safety				Memory safety [2014]				
(a) Verification	Time	N	# Z3	# Cl. (✓)	Time	N	# Z3	# Cl. (✓)	Time	N	# Z3	# Cl.	
Singly-linked lists													
concat	1.2	3	56	7 (4)	0.5	3	42	4 (1)	AF				
create	1.4	3	62	7 (1)	0.8	3	40	5 (1)	1.8	3	28	3	
delete	13.2	5	278	24 (11)	1.5	3	67	9 (1)	5.1	4	91	10	
delete-all	10.3	5	255	15 (8)	0.6	3	40	3 (1)	1.6	3	56	5	
filter	37.1	6	430	26 (16)	1.9	3	80	10 (3)	9.9	5	123	8	
insert-at	1.8	3	69	8 (1)	1.5	4	60	9 (1)	6.6	4	130	14	
insert	1.5	3	68	9 (2)	0.8	3	54	7 (1)	0.8	3	67	9	
merge	244.6	7	1429	36 (21)	10.7	6	260	13 (1)	AF				
reverse	19.7	5	289	13 (9)	2.7	5	114	5 (1)	5.0	6	246	4	
split	178.0	8	1079	33 (17)	5.9	5	146	11	7.0	5	97	9	
uf-find	43.3	8	590	25 (18)	3.1	8	144	4	6.1	10	266	8	
uf-union	178.8	7	1189	28 (13)	92.1	9	974	20 (5)	TO				
Sorted singly linked lists													
sorted-insert	3.9	3	85	13 (5)	1.2	3	56	8 (1)	6.8	3	63	8	
sorted-merge	400.3	8	1535	34 (21)	29.1	6	414	11 (2)	AF				
bubble-sort	90.5	11	904	23 (9)	2.1	5	61	6 (2)	2.3	4	34	3	
insertion-sort	1681.1	13	4601	48 (25)	100.2	11	1220	35 (5)	TO				
Doubly linked lists													
create	5.7	5	139	7 (3)	3.0	5	104	7 (2)	60.9	4	99	8	
delete	3.1	4	90	10 (5)	0.8	3	36	5 (2)	38.5	4	78	8	
insert-at	4.8	4	132	17 (8)	1.7	3	64	10 (3)	141.4	5	151	12	
Composite linked-list structures													
nested-flatten	564.6	17	3019	39 (24)	169.5	17	1810	22 (11)	AF				
nested-split	660.5	9	1144	27 (21)	6.8	4	143	9 (1)	AF				
overlaid-delete	188.5	6	1054	27 (4)	101.0	7	843	24 (2)	TO				
Restricted cyclic singly linked lists													
is-cycle	1.5	4	81	6	0.1	2	3	0	0.1	2	3	0	
last	4.6	6	143	8 (1)	1.6	4	56	5 (1)	10.2	8	155	5	
unchain	823.8	9	1986	50 (30)	0.1	2	3	0	0.1	2	3	0	
insert	91.2	5	208	18 (6)	9.6	5	101	11 (2)	405.1	5	127	10	
delete	8.2	4	138	16 (4)	5.1	4	95	10 (1)	55.4	3	81	13	
reverse	253.9	8	1202	24 (10)	20.2	8	319	13 (4)	84.1	7	421	12	
(b) Absence of a universal invariant													
					Description					Time	N	Z3	
shared-tail					See Example 5.8					3.6	2	42	
comb					See Section 7(b)					2	3	52	
(c) Bug finding					Bug description					Time	N	Z3	C.e. size
insert-at					Precondition is too weak (omitted $e \neq \text{null}$)					0.4	1	11	4
filter					Forgot a corner case where $\neg ok(h)$					3	1	21	4
insertion-sort					Typo: typed j instead of i					5	4	68	4
sorted-merge					Forgot to link the two segments					7.5	1	49	4

“[2014]” stands for results obtained using the analysis of Itzhaky et al. [2014]. N denotes the highest index for a developed frame F_i . “# Z3” denotes the number of calls to Z3. **AF** denotes “Abstraction Failure” of Itzhaky et al. [2014]. **TO** means timeout (> 1 hr). (a) Correct programs; “# Cl. (✓)” = number of (✓)-clauses in the inferred invariant, after the elimination of redundant clauses. (b) Correct programs for which there is no universal inductive invariant. (c) Incorrect programs; “C.e. size” is the maximal number of elements in a model that arises in the counterexample.

Table VI. Instrumentation Predicates That Were Manually Provided in the Previous Work [Itzhaky et al. 2014]

Name	Formula
$x\langle f \rangle y$	$\varphi_f(x, y)$ (see Table II)
$f.ls[x, y]$	$\forall \alpha, \beta. f^*(x, \alpha) \wedge f^*(\alpha, y) \wedge f^*(\beta, \alpha) \rightarrow (f^*(\beta, x) \vee f^*(x, \beta))$
$f.stable(h)$	$\forall \alpha. f^*(h, \alpha) \rightarrow alloc(\alpha)$
$f/b.rev[x, y]$	$\forall \alpha, \beta. \left(\begin{array}{l} \alpha \neq null \wedge \beta \neq null \wedge \\ f^*(x, \alpha) \wedge f^*(\alpha, y) \wedge f^*(x, \beta) \wedge f^*(\beta, y) \end{array} \right) \rightarrow (f^*(\alpha, \beta) \leftrightarrow b^*(\beta, \alpha))$
$f.sorted[x, y]$	$\forall \alpha, \beta. \left(\begin{array}{l} \alpha \neq null \wedge \beta \neq null \wedge \\ f^*(x, \alpha) \wedge f^*(\alpha, \beta) \wedge f^*(\beta, y) \end{array} \right) \rightarrow dle(\alpha, \beta)$

f and b denote pointer fields. dle is an uninterpreted predicate that denotes a total order on the data values; its semantics is enforced by an appropriate total-order background theory.

by a binary predicate together with the appropriate axioms. Using this predicate, we express sortedness, without directly considering the data.

In addition, we verified several procedures that manipulate multi-linked lists: `overlaid-delete()` takes an overlaid list and deletes a given element. (Overlaid lists use multiple pointer fields to index the same set of elements in different orders.) `nested-split()` moves all the elements not satisfying *ok* into a sublist. `flatten()` takes a nested list and flattens it by concatenating its sublists. We also verify the union-find algorithm. For example, for compressing `find()` operation, we prove that it maintains the reachability between every node and its root and preserves the elements.

Finally, we verified several procedures manipulating cyclic linked lists that contain not more than one cycle. `is-cycle()` checks if a given list is cyclic. `unchain()` breaks every *n*-edge of a given cyclic list. `delete()` deletes a given element from an input cyclic list. `insert()` insert a given element into a given sorted cyclic list while preserving the sortedness.

We compared our results to Itzhaky et al. [2014], where EA^R was used to verify properties of list-manipulating programs with PDR, using the human-supplied (universally quantified) abstraction predicate templates given in Table VI. The templates are parameterized for arbitrary pointer fields f and b and program variables x , y , and h . They have the following meanings:

- $x\langle f \rangle y$ — x is the immediate f -successor of y ;
- $f.ls[x, y]$ —the list segment from x to y is not shared by any heap location;
- $f.stable(h)$ — h is the head of a list containing only allocated objects;
- $f/b.rev[x, y]$ —field b is a back-pointer for field f in list segment from x to y ;
- $f.sorted[x, y]$ —the list segment from x to y is sorted.

Note that some of the predicates are far from trivial, and some are specific for one or two examples. Our technique was able to discover similar properties without the need for manually provided instrumentation predicates. We note that Itzhaky et al. [2014] can also establish certain functional correctness properties, but theirs are strictly weaker than ours. For example, they do not verify that a reversed list does not contain more elements than in its input list.

(b) *Verifying the Absence of Universal Invariants.* Our tool was also able to show that certain properties cannot be verified with a universal invariant. It proved that procedure `shared-tail()`, described in Example 5.8, does not have a universal invariant.

We applied our tool to procedure `comb()`, which for a given list `h` and every element `a` of `h` allocates a new element `b` and places a pointer `p` from `a` to `b`, hence resulting in a heap shaped like a comb. The tool discovered that it is not possible to use a universal invariant to prove that when `comb()` terminates there is no null-valued `p`-field in the input list.

(c) *Bug Finding.* We also ran our analysis on programs containing deliberate bugs. In all of the cases, the method was able to detect the bug and generate a concrete trace in which the safety or correctness properties are violated.

In our experience, the tool works well when the body of the loop is simple, even if the resulting invariants are not small. However, as the loop's body becomes more complicated, in particular containing conditional statements, the first-order formula describing the transition system can grow quite large, which has an adverse effect on scalability. This is most evident in the verification of `insertion-sort()`, where a nested loop was encoded using a single loop and some additional control statements.

8. RELATED WORK

Synthesizing quantified invariants has received significant attention. Several works have considered discovery of quantified predicates, for example, based on counterexamples [Das and Dill 2002] or by extension of predicate abstraction to support free variables [Flanagan and Qadeer 2002; Lahiri and Bryant 2007]. Our inferred invariants are comprised of universally quantified predicates, but unlike these approaches, our computation of the predicates is property directed and does *not* employ predicate abstraction. Additional works for generation of quantified invariants include using abstract domains of *quantified data automata* [Garg et al. 2013a, 2013b] or ones tailored to Presburger arithmetic with arrays [Dillig et al. 2010], instantiating quantifier templates [Björner et al. 2013; Srivastava and Gulwani 2009], or applying symbolic proof techniques [Hoder et al. 2011].

Other works aim to identify loop invariants *given* a set of predicates as candidate ingredients. Houdini [Flanagan and Leino 2001] is the first such algorithm of which we are aware. Santini [Thakur et al. 2013, 2015] is a recent algorithm that is based on full predicate abstraction. In the context of IC3, predicate abstraction was used in Birgmeier et al. [2014], Cimatti et al. [2014], and Itzhaky et al. [2014], the last of which specifically targeting shape analysis. The above require a predefined set of predicates and are therefore less automatic than our approach, since the diagrams provide an “on-the-fly” abstraction mechanism. An approach based on counterexample-guided refinement (CEGAR) for inferring predicates is shown in Podelski and Wies [2010]. They go beyond the usual state predicates by introducing unary predicates that range over heap locations. This allows us to infer quantified properties and permits lazy refinement to improve performance but does not provide completeness guarantees. In addition, the inferred invariants have quantifier-nesting depth of 1, whereas PDR^\forall allows for arbitrary nesting depth.

PDR has been shown to work extremely well in other domains, such as hardware verification [Bradley 2011; Eén et al. 2011]. Subsequently, it was generalized to software model checking for program models that use linear real arithmetic. The approach in Hoder and Björner [2012] uses conflict-based projection for linear real arithmetic, while Cimatti and Griggio [2012] employs a quantifier-elimination procedure for linear real arithmetic to provide an approximate pre-image operation. Finally, Komuravelli et al. [2014] uses model-based projection to extract under-approximations of pre-images. In contrast, our use of diagrams allows us to obtain a natural approximation that is precise for programs that can be verified using universal invariants.

The reduction we use into EPR creates a parametrized array-based system (where the range of the arrays are Booleans). A number of tools have been developed for general array-based systems. The SAFARI (SMT-Based Abstraction for Arrays with Interpolants) [Alberti et al. 2012] system is relevant. It is related to MCMT (Model Checker Modulo Theories) and Cubicle [Ghilardi and Ranise 2010a, 2010b; Conchon et al. 2012, 2013]. SAFARI uses symbolic preconditions to propagate symbolic states in the form of cubes that are conjunctions of literals over array constraints and uses interpolants to synthesize universal invariants. Our method for propagating and inductively generalizing diagrams differs by being based on PDR. More generally, our use of diagrams lazily produces finite-state abstractions of array-based systems. Lazy abstraction was applied in Henzinger et al. [2002] and shown to terminate when the abstraction structures satisfy the ascending-chain condition, for example, that every chain of abstractions is well founded. In Jhala and McMillan [2006], various abstraction domains are considered based on interpolation where termination of the interpolation process can be enforced by limiting the creation of atomic formulae to use only existing sub-terms. In contrast, our termination argument is based on well-quasi-orders.

The logic used by our implementation has limited capabilities to express properties of list segments that are not pointed to by variables [Itzhaky et al. 2014]. This is similar to the self-imposed limitations on expressibility used in a number of past approaches, including (a) canonical abstraction [Sagiv et al. 2002], (b) a prior method for applying predicate abstraction to linked lists [Manevich et al. 2005], (c) an abstraction method based on “must-paths” between nodes that are either pointed to by variables or are list-merge points [Lev-Ami et al. 2006], and (d) domains based on separation logic’s list-segment primitive [Distefano et al. 2006; Berdine et al. 2007] (i.e., “ $ls[x, y]$ ” asserts the existence of a possibly empty list segment running from the node pointed to by x to the node pointed to by y). Decision procedures have been used in previous work to compute the best transformer for individual statements that manipulate linked lists [Yorsh et al. 2004; Podelski and Wies 2010].

Several logics used for heap verification have decision procedures obtained by reduction to traditional logics. STRAND (STRucture AND Data) logic [Madhusudan and Qiu 2011; Madhusudan et al. 2011] has the ability to reason about heap data structures and the data they contain, using monadic second-order logic (MSO) on trees to define relations describing heaps. It has a decidable fragment inspired by EPR where universally quantified variables can be used only in so-called *elastic relation*. Essentially, elasticity is a generalization of transitive closure. STRAND is similar to the earlier PALE (Pointer Assertion Logic Engine) [Møller and Schwartzbach 2001], which also translates reachability properties to MSO and is based on *graph types* [Cook and Oppen 1975]. STRAND does not have an invariant inference mechanism, and, thus, it can be interesting to use STRAND as the logic \mathcal{L} in PDR^\forall . The logic CSL (Composite Structures Logic) [Bouajjani et al. 2009] has a similar flavor to STRAND, with similar sort restrictions on the syntax, but generalizes to handle doubly linked lists and allows size constraints on structures. Its decidability is obtained by proving a small model property and via a reduction to first-order logic.

A logic for reasoning about (cyclic) singly linked lists is proposed by Rakamaric et al. [2007]. The logic contains transitive closure of a single link function symbol and has a decision procedure based on custom inference rules. A logic for reasoning about list segments and data is given in Lahiri and Qadeer [2008]. The logic, LISBQ (Logic of Interpreted Sets and Bounded Quantification), provides a ternary primitive $\cdot \rightarrow \cdot \rightarrow \cdot$ that corresponds to heap paths through *three* nodes. This allows reasoning about pointer cycles. The decision procedure is based on a custom proof system, and its termination relies on stratification of the sorts—a semantic property of the formulae.

A previous work by the same authors [Lahiri and Qadeer 2006] proposes a translation to first-order logic but requires manual instantiations of quantifiers.

Separation logic [Reynolds 2002] uses inductive predicates conjoined with the $*$ -operator to describe unbounded heaps. It has been used as a basis for static shape analyses, for example, Distefano et al. [2006], and also has some decidable fragments, for example, Berdine et al. [2004]. The latter allows to describe heaps as a collection of separated list segments between program variables or existentially quantified ones. Our logic uses similar segments; however, it does not express separation explicitly but rather as a collection of properties that hold at all the heaps. TREX (Trees with Reachability EXpressions) [Wies et al. 2011] and GRIT (Graph Reachability and Inverted Trees) [Piskac et al. 2014a, 2014b] are essentially similar decidable fragments of separation logic. The former has a decision procedure based on reduction to first-order logic, and the latter was recently shown to be reducible to EPR. GRIT also allows to reason on numerical data stored in heap cells.

9. CONCLUSIONS

PDR^\forall is a combination of PDR/IC3 [Bradley 2011] with the model-theoretic notion of diagrams [Chang and Keisler 1990]. The latter provide PDR an aggressive strengthening scheme in which the structural properties of a bad state are abstracted “on-the-fly” by a formula describing all of its possible extensions, which are then blocked together within the same iteration of PDR’s main refinement loop. This obviates the need for user-supplied abstraction predicates. This form of automation is particularly important when one tries to verify tricky programs, for example, programs that manipulate unbounded data structures, against a variety (of possibly changing) specifications. Indeed, our implementation successfully analyzed multiple specifications of tricky list-manipulating programs, discovered counterexamples, and, uniquely to our approach, showed that certain programs cannot be proven correct using a universal invariant. Interestingly, we noticed that sometimes the tool had to work harder to verify simple properties than when it was asked to verify complicated ones. In particular, verifying partial correctness properties was done faster when verified together with memory safety than without. In hindsight, this might not be surprising due to the property guided nature of the analysis.

We are very pleased with the simplicity of our approach and believe that the notion of diagram-based abstractions is particularly useful for the verification of programs that manipulate unbounded state. Recent work [Frumkin et al. 2017] has extended PDR^\forall to interprocedural analysis, where procedure summaries [Reps et al. 1995] are inferred instead of invariants. The interprocedural version was used to verify correct use of iterators in Java programs. In the future, we plan to apply it in other contexts too, for example, for the verification of network programs [ONF 2016].

APPENDIX

A. THE INVARIANT OBTAINED FOR SPLIT()

$$\begin{aligned}
I &= \bigwedge_{i=0}^{32} L_i \quad \text{where} \\
L_0 &= (j \neq i) \vee (i = h) \\
L_1 &= \forall \alpha. (ok(\alpha) \vee n^*(\alpha, k) \vee (g = \underline{h}) \vee \underline{n}^*(j, \alpha) \vee \neg \underline{n}^*(g, \alpha)) \\
L_2 &= ((g = null) \vee \neg ok(g)) \\
L_3 &= (\underline{n}^*(k, j) \vee \neg n^*(j, k)) \\
L_4 &= ((h = i) \vee \underline{n}^*(j, i) \vee (null = i)) \\
L_5 &= (\underline{n}^*(j, i) \vee (j = null) \vee (null = i)) \\
L_6 &= \forall \alpha \beta. (n^*(\underline{h}, \alpha) \vee \neg \underline{n}^*(\underline{h}, \alpha) \vee \neg n^*(\underline{h}, \beta) \vee \neg n^*(\alpha, \beta)) \\
L_7 &= \forall \alpha. (\neg \underline{n}^*(\underline{h}, \alpha) \vee ok(\alpha) \vee (g = null) \vee \underline{n}^*(g, \alpha)) \\
L_8 &= \forall \alpha \beta. (\underline{n}^*(\alpha, \underline{h}) \vee \neg \underline{n}^*(\alpha, \beta) \vee n^*(\alpha, \beta) \vee \underline{n}^*(\alpha, k)) \\
L_9 &= (\neg n^*(\underline{h}, h) \vee \underline{n}^*(h, \underline{h})) \\
L_{10} &= \forall \alpha. (\underline{n}^*(i, \alpha) \vee \neg \underline{n}^*(j, \alpha) \vee (\alpha = j) \vee \underline{n}^*(j, k)) \\
L_{11} &= (ok(j) \vee (j = null)) \\
L_{12} &= (n^*(h, j) \vee (j = null)) \\
L_{13} &= \forall \alpha. (n^*(g, \alpha) \vee ok(\alpha) \vee \neg \underline{n}^*(\underline{h}, \alpha) \vee (g = k) \vee \underline{n}^*(k, \alpha)) \\
L_{14} &= ((\underline{h} = g) \vee (\underline{h} = h)) \\
L_{15} &= \forall \alpha \beta. (\neg n^*(\alpha, \beta) \vee \underline{n}^*(\alpha, \beta)) \\
L_{16} &= \underline{n}^*(g, k) \\
L_{17} &= \forall \alpha. (\neg \underline{n}^*(\underline{h}, \alpha) \vee n^*(\underline{h}, \alpha) \vee \underline{n}^*(g, \alpha)) \\
L_{18} &= \forall \alpha \beta. (\underline{n}^*(k, \beta) \vee \neg \underline{n}^*(\alpha, \beta) \vee n^*(\alpha, \beta) \vee \neg ok(\beta) \vee \neg \underline{n}^*(\underline{h}, \alpha) \vee n^*(\alpha, k)) \\
L_{19} &= \forall \alpha. (n^*(\underline{h}, \alpha) \vee \neg \underline{n}^*(\underline{h}, \alpha) \vee \underline{n}^*(h, \alpha)) \\
L_{20} &= ((null = i) \vee \underline{n}^*(k, i) \vee (g = null)) \\
L_{21} &= \forall \alpha. ((h = i) \vee \underline{n}^*(j, \alpha) \vee (k \neq null) \vee ok(\alpha) \vee \neg \underline{n}^*(h, \alpha)) \\
L_{22} &= \forall \alpha. (n^*(\alpha, k) \vee \neg ok(\underline{h}) \vee \underline{n}^*(\alpha, g) \vee \underline{n}^*(j, \underline{h}) \vee \underline{n}^*(\alpha, j) \vee \neg \underline{n}^*(\alpha, k)) \\
L_{23} &= \forall \alpha \beta. (\underline{n}^*(j, \alpha) \vee (\alpha = k) \vee \neg \underline{n}^*(k, \alpha) \vee \underline{n}^*(k, \beta) \vee \neg \underline{n}^*(\underline{h}, \beta) \vee n^*(\beta, \alpha) \vee n^*(\beta, k)) \\
L_{24} &= \forall \alpha. ((\alpha = k) \vee \underline{n}^*(i, \alpha) \vee \neg \underline{n}^*(k, \alpha) \vee \underline{n}^*(k, j)) \\
L_{25} &= ((h = null) \vee \underline{n}^*(h, h)) \\
L_{26} &= (\underline{n}^*(\underline{h}, g) \vee (g = null)) \\
L_{27} &= \forall \alpha. (\neg \underline{n}^*(j, \alpha) \vee \neg \underline{n}^*(\alpha, k) \vee n^*(\alpha, k) \vee \underline{n}^*(\alpha, j)) \\
L_{28} &= \forall \alpha. (\neg n^*(g, \alpha) \vee \underline{n}^*(\alpha, g) \vee \neg ok(\alpha)) \\
L_{29} &= ((\underline{h} = h) \vee (k \neq h)) \\
L_{30} &= \forall \alpha \beta. (\neg \underline{n}^*(\alpha, \beta) \vee ok(\beta) \vee \underline{n}^*(j, \beta) \vee (\alpha = \beta) \vee \neg \underline{n}^*(k, j) \vee \neg \underline{n}^*(k, \alpha)) \\
L_{31} &= (n^*(h, i) \vee (i = null)) \\
L_{32} &= (\underline{n}^*(j, g) \vee (i \neq g) \vee (g = null))
\end{aligned}$$

Fig. 7. The inferred invariant for split. It contains 33 clauses, of which 17 are universal.

ACKNOWLEDGMENTS

We thank Mooly Sagiv and the reviewers of this article and of the preliminary conference version [Karbyshev et al. 2015] for their helpful comments.

REFERENCES

- Aws Albarghouthi, Josh Berdine, Byron Cook, and Zachary Kincaid. 2015. Spatial interpolants. *CoRR* abs/1501.04100 (2015). Retrieved from <http://arxiv.org/abs/1501.04100>
- Francesco Alberti, Roberto Bruttomesso, Silvio Ghilardi, Silvio Ranise, and Natasha Sharygina. 2012. SAFARI: SMT-based abstraction for arrays with interpolants. In *Proceedings of 24th International Conference on Computer Aided Verification (CAV'12)*, Vol. 7358. Springer. DOI: http://dx.doi.org/10.1007/978-3-642-31424-7_49
- Clark Barrett, Aaron Stump, and Cesare Tinelli. 2010. The Satisfiability Modulo Theories Library (SMT-LIB). (2010). Retrieved from <http://www.smt-lib.org>.
- Josh Berdine, Cristiano Calcagno, Byron Cook, Dino Distefano, Peter W. O'Hearn, Thomas Wies, and Hongseok Yang. 2007. Shape analysis for composite data structures. In *Proceedings of 19th International Conference on Computer Aided Verification (CAV'07)*, Vol. 4590. Springer, 178–192. DOI: http://dx.doi.org/10.1007/978-3-540-73368-3_22
- Josh Berdine, Cristiano Calcagno, and Peter W. O'Hearn. 2004. A decidable fragment of separation logic. In *Proceedings of the 24th International Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS'04)*. Springer-Verlag, Berlin, 97–109. DOI: http://dx.doi.org/10.1007/978-3-540-30538-5_9
- Armin Biere, Alessandro Cimatti, Edmund M. Clarke, Ofer Strichman, and Yunshan Zhu. 2003. Bounded model checking. *Adv. Comput.* 58 (2003), 118–149.
- Armin Biere, Alessandro Cimatti, Edmund M. Clarke, and Yunshan Zhu. 1999. Symbolic model checking without BDDs. In *Proceedings of 5th International Conference on Tools and Algorithms for Construction and Analysis of Systems (TACAS'99)*. Springer, 193–207. DOI: http://dx.doi.org/10.1007/3-540-49059-0_14
- Johannes Birgmeier, Aaron R. Bradley, and Georg Weissenbacher. 2014. Counterexample to induction-guided abstraction-refinement (CTIGAR). In *Proceedings of 26th International Conference on Computer Aided Verification (CAV'14)*, Vol. 8559. Springer, 831–848. DOI: http://dx.doi.org/10.1007/978-3-319-08867-9_55
- Nikolaj Bjørner, Kenneth L. McMillan, and Andrey Rybalchenko. 2013. On solving universally quantified horn clauses. In *Proceedings of 20th International Static Analysis Symposium (SAS'13)*, Vol. 7935. Springer, 105–125. DOI: http://dx.doi.org/10.1007/978-3-642-38856-9_8
- Egon Börger, Erich Grädel, and Yuri Gurevich. 2008. *The Classical Decision Problem*. Springer-Verlag.
- Ahmed Bouajjani, Cezara Dragoi, Constantin Enea, and Mihaela Sighireanu. 2009. A logic-based framework for reasoning about composite data structures. In *Proceedings of 20th International Conference on Concurrency Theory (CONCUR'09)*, Vol. 5710. Springer, 178–195. DOI: http://dx.doi.org/10.1007/978-3-642-04081-8_13
- Aaron R. Bradley. 2011. SAT-based model checking without unrolling. In *Proceedings of the 12th International Conference on Verification, Model Checking, and Abstract Interpretation (VMCAI'11)*, Vol. 6538. Springer, 70–87. DOI: http://dx.doi.org/10.1007/978-3-642-18275-4_7
- Chen C. Chang and Howard J. Keisler. 1990. *Model Theory*. Elsevier Science.
- Alessandro Cimatti and Alberto Griggio. 2012. Software model checking via IC3. In *Proceedings of the 24th International Conference on Computer Aided Verification (CAV'12)*, Vol. 7358. Springer, 277–293. DOI: http://dx.doi.org/10.1007/978-3-642-31424-7_23
- Alessandro Cimatti, Alberto Griggio, Sergio Mover, and Stefano Tonetta. 2014. IC3 modulo theories via implicit predicate abstraction. In *Proceedings of the 20th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'14)*, Vol. 8413. Springer, 46–61. DOI: http://dx.doi.org/10.1007/978-3-642-54862-8_4
- Edmund Clarke, Daniel Kroening, and Karen Yorav. 2003. Behavioral consistency of C and verilog programs using bounded model checking. In *Proceedings of the 40th Annual Design Automation Conference (DAC'03)*. ACM, New York, NY, 368–371. DOI: <http://dx.doi.org/10.1145/775832.775928>
- Sylvain Conchon, Amit Goel, Sava Krstic, Alain Mebsout, and Fatiha Zaïdi. 2012. Cubicle: A parallel smt-based model checker for parameterized systems—tool paper. In *Proceedings of the 24th International Conference on Computer Aided Verification (CAV'12)*, Vol. 7358. Springer, 718–724. DOI: http://dx.doi.org/10.1007/978-3-642-31424-7_55
- Sylvain Conchon, Amit Goel, Sava Krstic, Alain Mebsout, and Fatiha Zaïdi. 2013. Invariants for finite instances and beyond. In *Proceedings of Conference on Formal Methods in Computer-Aided Design (FMCAD'13)*. IEEE, 61–68.

- Stephen A. Cook and Derek C. Oppen. 1975. An assertion language for data structures. In *Proceedings of the 2Nd ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages (POPL'75)*. ACM, New York, NY, 160–166. DOI: <http://dx.doi.org/10.1145/512976.512993>
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. 2009. *Introduction to Algorithms* (3rd ed.). The MIT Press.
- Patrick Cousot and Radhia Cousot. 1977. Abstract interpretation: A unified lattice model for static analysis of programs by construction or approximation of fixpoints. In *Conference Record of the Fourth ACM Symposium on Principles of Programming Languages (POPL'77)*. ACM, 238–252. DOI: <http://dx.doi.org/10.1145/512950.512973>
- Satyaki Das and David L. Dill. 2002. Counter-example based predicate discovery in predicate abstraction. In *Proceedings of the 4th International Conference on Formal Methods in Computer-Aided Design (FMCAD'02)*, Vol. 2517. Springer, 19–32. DOI: http://dx.doi.org/10.1007/3-540-36126-X_2
- Leonardo Mendonça de Moura and Nikolaj Bjørner. 2008. Z3: an efficient SMT solver. In *Proceedings of the 14th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'08)*, Vol. 4963. Springer, 337–340. DOI: http://dx.doi.org/10.1007/978-3-540-78800-3_24
- Isil Dillig, Thomas Dillig, and Alex Aiken. 2010. Symbolic heap abstraction with demand-driven axiomatization of memory invariants. In *ACM SIGPLAN Conference on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA'10)*. ACM, 397–410. DOI: <http://dx.doi.org/10.1145/1869459.1869493>
- Dino Distefano, Peter W. O'Hearn, and Hongseok Yang. 2006. A local shape analysis based on separation logic. In *Proceedings of the 12th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'06)*, Vol. 3920. Springer, 287–302. DOI: http://dx.doi.org/10.1007/11691372_19
- Niklas Eén, Alan Mishchenko, and Robert K. Brayton. 2011. Efficient implementation of property directed reachability. In *Proceedings of International Conference on Formal Methods in Computer-Aided Design (FMCAD'11)*. FMCAD Inc., 125–134.
- Cormac Flanagan and K. Rustan M. Leino. 2001. Houdini, an annotation assistant for ESC/java. In *Proceedings of Symposium of Formal Methods Europe (FME'01)*, Vol. 2021. Springer, 500–517. DOI: http://dx.doi.org/10.1007/3-540-45251-6_29
- Cormac Flanagan and Shaz Qadeer. 2002. Predicate abstraction for software verification. In *Conference Record of the 29th SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'02)*. ACM, 191–202. DOI: <http://dx.doi.org/10.1145/503272.503291>
- Asya Frumkin, Yotam M. Y. Feldman, Ondřej Lhoták, Oded Padon, Mooly Sagiv, and Sharon Shoham. 2017. Property directed reachability for proving absence of concurrent modification errors. In *Proceedings of the 18th International Conference on Verification, Model Checking, and Abstract Interpretation (VMCAI'17)*. Springer. To appear.
- Pranav Garg, Christof Löding, P. Madhusudan, and Daniel Neider. 2013a. Learning universally quantified invariants of linear data structures. In *Proceedings of the 25th International Conference on Computer Aided Verification (CAV'13)*, Vol. 8044. Springer, 813–829. DOI: http://dx.doi.org/10.1007/978-3-642-39799-8_57
- Pranav Garg, P. Madhusudan, and Gennaro Parlato. 2013b. Quantified data automata on skinny trees: An abstract domain for lists. In *Proceedings of the 20th International Symposium on Static Analysis (SAS'13)*, Vol. 7935. Springer, 172–193. DOI: http://dx.doi.org/10.1007/978-3-642-38856-9_11
- Silvio Ghilardi and Silvio Ranise. 2010a. Backward reachability of array-based systems by SMT solving: Termination and invariant synthesis. *Logic. Methods Comput. Sci.* 6, 4 (2010). DOI: [http://dx.doi.org/10.2168/LMCS-6\(4:10\)2010](http://dx.doi.org/10.2168/LMCS-6(4:10)2010)
- Silvio Ghilardi and Silvio Ranise. 2010b. MCMT: A model checker modulo theories. In *Proceedings of the 5th International Joint Conference on Automated Reasoning (IJCAR'10)*, Vol. 6173. Springer, 22–29. DOI: http://dx.doi.org/10.1007/978-3-642-14203-1_3
- Thomas A. Henzinger, Ranjit Jhala, Rupak Majumdar, and Grégoire Sutre. 2002. Lazy abstraction. In *Conference Record of POPL 2002: The 29th SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. ACM, 58–70. DOI: <http://dx.doi.org/10.1145/503272.503279>
- Krystof Hoder and Nikolaj Bjørner. 2012. Generalized property directed reachability. In *Proceedings of the 15th International Conference on Theory and Applications of Satisfiability Testing (SAT'12)*, Vol. 7317. Springer, 157–171. DOI: http://dx.doi.org/10.1007/978-3-642-31612-8_13
- Krystof Hoder, Laura Kovács, and Andrei Voronkov. 2011. Invariant generation in vampire. In *Proceedings of the 17th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'11)*, Vol. 6605. Springer, 60–64. DOI: http://dx.doi.org/10.1007/978-3-642-19835-9_7
- Shachar Itzhaky, Anindya Banerjee, Neil Immerman, Aleksandar Nanovski, and Mooly Sagiv. 2013. Effectively-propositional reasoning about reachability in linked data structures. In *Proceedings of the*

- 25th International Conference on Computer Aided Verification (CAV'13)*, Vol. 8044. Springer, 756–772. DOI : http://dx.doi.org/10.1007/978-3-642-39799-8_53
- Shachar Itzhaky, Nikolaj Bjørner, Thomas W. Reps, Mooly Sagiv, and Aditya V. Thakur. 2014. Property-directed shape analysis. In *Proceedings of the 26th International Conference on Computer Aided Verification (CAV'14)*, Vol. 8559. Springer, 35–51. DOI : http://dx.doi.org/10.1007/978-3-319-08867-9_3
- Ranjit Jhala and Kenneth L. McMillan. 2006. A practical and complete approach to predicate refinement. In *Proceedings of the 12th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'06)*, Vol. 3920. Springer, 459–473. DOI : http://dx.doi.org/10.1007/11691372_33
- Aleksandr Karbyshev, Nikolaj Bjørner, Shachar Itzhaky, Noam Rinetzkzy, and Sharon Shoham. 2015. Property-directed inference of universal invariants or proving their absence. In *Proceedings of the 27th International Conference on Computer Aided Verification (CAV'15)*, Vol. 9206. Springer, 583–602. DOI : http://dx.doi.org/10.1007/978-3-319-21690-4_40
- Anvesh Komuravelli, Arie Gurfinkel, and Sagar Chaki. 2014. SMT-based model checking for recursive programs. In *Proceedings of the 26th International Conference on Computer Aided Verification (CAV'14)*, Vol. 8559. Springer, 17–34. DOI : http://dx.doi.org/10.1007/978-3-319-08867-9_2
- Shuvendu K. Lahiri and Randal E. Bryant. 2007. Predicate abstraction with indexed predicates. *ACM Trans. Comput. Logic* 9, 1 (2007).
- Shuvendu K. Lahiri and Shaz Qadeer. 2006. Verifying properties of well-founded linked lists. In *Proceedings of the 33rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'06)*. ACM, New York, NY, 115–126. DOI : <http://dx.doi.org/10.1145/1111037.1111048>
- Shuvendu K. Lahiri and Shaz Qadeer. 2008. Back to the future: Revisiting precise program verification using SMT solvers. In *Proceedings of the 35th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'08)*. ACM, 171–182. DOI : <http://dx.doi.org/10.1145/1328438.1328461>
- Tal Lev-Ami, Neil Immerman, and Shmuel Sagiv. 2006. Abstraction for shape analysis with fast and precise transformers. In *Proceedings of the 18th International Conference on Computer Aided Verification (CAV'06)*, Vol. 4144. Springer, 547–561. DOI : http://dx.doi.org/10.1007/11817963_49
- Harry R. Lewis. 1980. Complexity results for classes of quantificational formulas. *J. Comput. Syst. Sci.* 21, 3 (1980), 317–353.
- P. Madhusudan, Gennaro Parlato, and Xiaokang Qiu. 2011. Decidable logics combining heap structures and data. In *Proceedings of the 38th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'11)*. ACM, New York, NY, 611–622. DOI : <http://dx.doi.org/10.1145/1926385.1926455>
- P. Madhusudan and Xiaokang Qiu. 2011. Efficient decision procedures for heaps using STRAND. In *Proceedings of the 18th International Symposium on Static Analysis (SAS'11)*, Vol. 6887. Springer, 43–59. DOI : http://dx.doi.org/10.1007/978-3-642-23702-7_8
- Roman Manevich, Eran Yahav, Ganesan Ramalingam, and Shmuel Sagiv. 2005. Predicate abstraction and canonical abstraction for singly-linked lists. In *Proceedings of the 6th International Conference on Verification, Model Checking, and Abstract Interpretation (VMCAI'05)*, Vol. 3385. Springer, 181–198. DOI : http://dx.doi.org/10.1007/978-3-540-30579-8_13
- Anders Møller and Michael I. Schwartzbach. 2001. The pointer assertion logic engine. In *Proceedings of the 2001 ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'01)*. ACM, New York, NY, 221–231. DOI : <http://dx.doi.org/10.1145/378795.378851>
- ONF. 2016. (2016). The Open Networking Foundation. Retrieved from <http://opennetworking.org>.
- Oded Padon, Neil Immerman, Sharon Shoham, Aleksandr Karbyshev, and Mooly Sagiv. 2016. Decidability of inferring inductive invariants. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'16)*. ACM, 217–231. DOI : <http://dx.doi.org/10.1145/2837614.2837640>
- Ruzica Piskac, Thomas Wies, and Damien Zufferey. 2014a. Automating separation logic with trees and data. In *Proceedings of 26th International Conference on Computer Aided Verification (CAV'14)*, Armin Biere and Roderick Bloem (Eds.), Vol. 8559. Springer, 711–728. DOI : http://dx.doi.org/10.1007/978-3-319-08867-9_47
- Ruzica Piskac, Thomas Wies, and Damien Zufferey. 2014b. GRASShopper - Complete heap verification with mixed specifications. In *Proceedings of the 20th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS 2014)*, Vol. 8413. Springer, 124–139. DOI : http://dx.doi.org/10.1007/978-3-642-54862-8_9
- Andreas Podelski and Thomas Wies. 2010. Counterexample-guided focus. In *Proceedings of the 37th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'10)*. ACM, 249–260. DOI : <http://dx.doi.org/10.1145/1706299.1706330>
- Zvonimir Rakamaric, Jesse D. Bingham, and Alan J. Hu. 2007. An inference-rule-based decision procedure for verification of heap-manipulating programs with mutable data and cyclic data structures. In *Proceedings*

- of the 8th International Conference on Verification, Model Checking, and Abstract Interpretation (VMCAI 2007), Vol. 4349. Springer, 106–121. DOI: http://dx.doi.org/10.1007/978-3-540-69738-1_8
- Thomas W. Reps, Susan Horwitz, and Shmuel Sagiv. 1995. Precise interprocedural dataflow analysis via graph reachability. In *Conference Record of the 22nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL95)*. ACM Press, 49–61. DOI: <http://dx.doi.org/10.1145/199448.199462>
- John C. Reynolds. 2002. Separation logic: A logic for shared mutable data structures. In *Proceedings of the 17th IEEE Symposium on Logic in Computer Science (LICS'02)*. IEEE Computer Society, 55–74. DOI: <http://dx.doi.org/10.1109/LICS.2002.1029817>
- Shmuel Sagiv, Thomas W. Reps, and Reinhard Wilhelm. 2002. Parametric shape analysis via 3-valued logic. *ACM Trans. Program. Lang. Syst.* 24, 3 (2002), 217–298. DOI: <http://dx.doi.org/10.1145/514188.514190>
- Saurabh Srivastava and Sumit Gulwani. 2009. Program verification using templates over predicate abstraction. In *Proceedings of the 2009 ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'09)*. ACM, 223–234. DOI: <http://dx.doi.org/10.1145/1542476.1542501>
- A. Thakur, A. Lal, J. Lim, and T. Reps. 2013. *PostHat and All That: Attaining Most-Precise Inductive Invariants*. TR-1790. Comp. Sci. Dept., Univ. of Wisconsin—Madison, Madison, WI.
- Aditya V. Thakur, Akash Lal, Junghee Lim, and Thomas W. Reps. 2015. Posthat and all that: Automating abstract interpretation. *Electr. Notes Theor. Comput. Sci.* 311 (2015), 15–32. DOI: <http://dx.doi.org/10.1016/j.entcs.2015.02.003>
- Thomas Wies, Marco Muñoz, and Viktor Kuncak. 2011. An efficient decision procedure for imperative tree data structures. In *Proceedings of the 23rd International Conference on Automated Deduction (CADE-23)*, Vol. 6803. Springer, 476–491. DOI: http://dx.doi.org/10.1007/978-3-642-22438-6_36
- Glynn Winskel. 1993. *The Formal Semantics of Programming Languages: An Introduction*. MIT Press, Cambridge, MA.
- Greta Yorsh, Thomas W. Reps, and Shmuel Sagiv. 2004. Symbolically computing most-precise abstract operations for shape analysis. In *Proceedings of the 10th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'04)*, Vol. 2988. Springer, 530–545. DOI: http://dx.doi.org/10.1007/978-3-540-24730-2_39

Received January 2016; revised November 2016; accepted December 2016